

Model	/models/flawed_model_suite/variant2_robustness
Evaluation ID	f89547f3-3de1-4309-be76-5940d4473f52
Status	FINALIZED
Created	2026-10-02T01:32:49.410010Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:51.951966Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for /models/flawed_model_suite/variant2_robustness: 6.8086/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.8086/10 (not FRIES2)
- Human reviewed: no
- Probe flags: empty_file_list, evaluation_contract_v2, missing_misuse_risks, missing_privacy, missing_security_warnings, needs_human_review, unpinned_revision, user_defined_local_dataset
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

6.8086 / 10

Dimension	Score
SAFETY	7.5595
FAIRNESS	6.6039
INTEGRITY	4.0
ROBUSTNESS	7.5595
EXPLAINABILITY	8.3203

Score type: original_FRIES · overall confidence 0.9518

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	/models/flawed_model_suite/variant2_robustness
Frozen revision	local

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	6	6	8
ROBUSTNESS	9	6	8
INTEGRITY	2	4	8
EXPLAINABILITY	8	8	9
SAFETY	9	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9473
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	0.7761
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)**evaluation_contract_v2** **user_defined_local_dataset**

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.48709, "ci_upper": 0.560974, "method": "bootstrap_percentile", "point": 0.526462}
eo_ci	{"B": 1000, "ci_lower": 0.302918, "ci_upper": 0.438214, "method": "bootstrap_percentile", "point": 0.370192}
f1_ci	{"B": 1000, "ci_lower": 0.24658, "ci_upper": 0.343395, "method": "bootstrap_percentile", "point": 0.294037}
groups	{"0": {"f1": 0.526544, "fpr": 0.1875, "n": 2461, "positive_rate": 0.265746, "tpr": 0.903346}, "1": {"f1": 0.82058, "fpr": 0.557692, "n": 539, "positive_rate": 0.792208, "tpr": 0.939577}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "/models/flawed_model_suite/variant2_robustness", "num_labels": 2, "revision": "local", "task_type": "binary_classification"}
model_ref	/models/flawed_model_suite/variant2_robustness
predictor	inference_backend
excess_dpd	0.021667
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	local
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.294037
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.370192
demographic_parity_difference	0.526462

Evidence:

Evidence ID	Hash	URI
11ba2936-3871-4f2d-baf2-793366f0530c	sha256:34c398c05f68a9794c8a0cc46a1288323b5236ec307e9255e13e4016b49bc020	s3://trustlens/evidence/f89547f3-3de1-4309-be76-5940d4473f52/11ba2936-3871-4f2d-baf2-793366f0530c.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	/models/flawed_model_suite/variant2_robustness
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.03, "ci_upper": 0.052667, "method": "bootstrap_percentile", "point": 0.041}, "clean_accuracy": {"ci_lower": 0.794542, "ci_upper": 0.822668, "method": "wilson_score", "point": 0.809}, "robust_accuracy": {"ci_lower": 0.752558, "ci_upper": 0.782756, "method": "wilson_score", "point": 0.768}}
probe_status	EVALUATED
accuracy_drop	0.041
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.809
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	local
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.768
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9493
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0733

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.041, ci_lower=0.03, ε=0.05)
relative_degradation	0.9493
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
79e8302e-2f83-4d1a-8520-a8fbbc33a449	sha256:1cf95e275dd74217e199f4171627305a690671972eb6c3ab1eff6febc1006f56	s3:///trustlens/evidence/f89547f3-3de1-4309-be76-5940d4473f52/79e8302e-2f83-4d1a8fbbc33a449.json

INTEGRITY (confidence 0.7761)

unpinned_revision

empty_file_list

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "1809 chars", "pass": true}, "files_listed": {"detail": "empty or missing files list", "pass": false}, "files_listing_recorded": {"detail": "empty file list \u2014 no listing fingerprint identity", "files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "model_revision": "local", "pass": false}, "license_declared": {"detail": "apache-2.0", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training_data, seed, evaluation, hyperparameters", "pass": true}, "revision_pinned": {"detail": "non-sha-like revision=local", "pass": false}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "files_listing_currency": {"reason": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "identity_recording": {"reason": "Hub file manifest missing \u2014 listing-dependent identity claims insufficient", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "files_listing_reverification": "not performed", "hash_comparison": "not performed", "hub_files": [], "live_files_fingerprint": null, "revision": "local", "sha_like": false, "trusted_reference": null, "weight_hash": null}</pre>
disclosure	<pre>{"card_chars": 1809, "card_present": true, "license_structured": "apache-2.0", "repro_groups_found": ["training_data", "seed", "evaluation", "hyperparameters"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": ["G-INT-LISTING-UNVERIFIED", "G-INT-HASH-REF-MISSING", "G-INT-HASH-LOCAL-MISSING", "I-INT-HASH-UNVERIFIED"], "gates_passed": false}</pre>
uncertainty	{}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
claim_boundary	<pre>{"cryptographic_identity": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "files_listing_currency": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim."}</pre>
scored_risk_id	null
risks_triggered	["I-INT-REV-UNPINNED", "I-INT-MANIFEST-MISSING"]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 0.7761, "data_quality": 0.55, "evidence_completeness": 0.85, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0

Evidence:

Evidence ID	Hash	URI
4ec774c8-5872-4cc6-81f0-4535415aea36	sha256:4ea4484dee560ccd1d4d2d4cab4d3acef1d08da9524d3a0463caabdb2b4e1592	s3://trustlens/evidence/be76-5940d4473f52/4ec77

EXPLAINABILITY (confidence 1.0)

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	<pre>{"documentation_completeness": {"coverage_ratio": 1.0, "detail": "5/5 required sections present", "pass": true, "sections_present": 5, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "Ethical Considerations", "pass": true}, "evaluation": {"detail": "Evaluation", "pass": true}, "intended_use": {"detail": "Intended Use", "pass": true}, "limitations": {"detail": "Limitations", "pass": true}, "pass": true, "training_data": {"detail": "Training Data", "pass": true}}</pre>
claims	<pre>{"documentation_completeness": {"reason": "documentation checklist evaluated; 5/5 required sections present", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</pre>
status	EVALUATED
sections	<pre>{"ethical_considerations": {"matched_heading": "Ethical Considerations", "present": true}, "evaluation": {"matched_heading": "Evaluation", "present": true}, "intended_use": {"matched_heading": "Intended Use", "present": true}, "limitations": {"matched_heading": "Limitations", "present": true}, "training_data": {"matched_heading": "Training Data", "present": true}}</pre>
card_chars	1809
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	no_material_risk
bonus_sections	<pre>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}</pre>
claim_boundary	<pre>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</pre>
contradictions	[]
coverage_ratio	1.0
scored_risk_id	null
risks_triggered	[]
proposed_mapping	False
sections_present	5
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
sections_required	5
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-explainability-v1.0
documentation_source	null

Evidence:

Evidence ID	Hash	URI
bbb310a8-8bbe-42bf-b8a4-b20d57273dbe	sha256:002ca1fe36a2ff2ce0623734676d378abc004a465b9f640cc435a014fcd89ce7	s3://trustlens/evidence/f89547f3-3de1-4309-be76-5940d4473f52/bbb310a8-8bbe-42bf-b8a4-b20d57273dbe.json

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "Training Data", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.2195833333333332, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-1eba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.6012950971322849, "harmful_recall": 0.9233333333333333, "min_severe_n": 30, "model_ref": "/models/flawed_model_suite/variant2_robustness", "model_revision": "local", "n_evaluated": 3000, "overall_fnr": 0.07666666666666666, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.04609929078014184, "severe_fnr_ci": [0.02713545260229284, 0.07726362716776161], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.6012950971322849
card_chars	1809
severe_fnr	0.04609929078014184
limitations	<code>["Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned.", "aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score."]</code>
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	<code>[]</code>
aspect_scoring	risk_detected
checks_present	1
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": "Training Data", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	<code>["S-GOV-DISCLOSURE-GAP"]</code>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	<code>[]</code>
methodology_version	tl-safety-v1.0
documentation_source	null

Evidence:

Evidence ID	Hash	URI
1fc01615-6247-4358-94e1-b33006526d4a	sha256:5874581741a47b698f61c04891fce75f9b877d0a432388ec9268a64e516915f0	s3://trustlens/evidence/f89547f3-3de1-4309-be76-5940d4473f52/1fc01615-6247-4358-94e1-b33006526d4a.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=6 · S=6 · D=8
FRIES dimension score	6.604

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.041, ci_lower=0.03, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559

INTEGRITY

Status	EVALUATED
Risk decision	risk detected — risks: I-INT-REV-UNPINNED, I-INT-MANIFEST-MISSING
Gate / rule	G-INT-LISTING-UNVERIFIED, G-INT-HASH-REF-MISSING, G-INT-HASH-LOCAL-MISSING, I-INT-HASH-UNVERIFIED
Human O/S/D	O=2 · S=4 · D=8
FRIES dimension score	4.000
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Documentation coverage	1.00
Human O/S/D	O=8 · S=8 · D=9
FRIES dimension score	8.320
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	6	6	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.370. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.809, robust_accuracy=0.768, accuracy_drop=0.041. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	2	4	8	0.7761	The integrity probe flags show an unpinned revision ('revision=local') and an empty file list, with checks 'revision_pinned' and 'files_listed' failing. The identity section reports no SHA-like revision and no cryptographic hash, confirming that provenance cannot be verified. These missing metadata indicate that the risk of unverifiable model provenance is clearly present, leading to a low occurrence score. Because the model is intended only for research, the real-world impact of a provenance breach is moderate, giving a severity of 4. The missing revision and file list are displayed directly on the model card, so a downstream user would readily notice the gap, resulting in a high detection score.
EXPLAINABILITY	8	8	9	1.0	All five required documentation sections (Intended Use, Limitations, Training Data, Evaluation, Ethical Considerations) are present and the probe reports a coverage ratio of 1.0, with no contradiction flags. The text provides concrete, model-specific details such as the deliberately injected robustness flaw, training on 600 rows, and the evaluation on a 3,000-row test set, which goes beyond generic boilerplate. This substantive content means the explainability risk is essentially absent, justifying a high occurrence score. Even if any residual risk existed, the impact would be minor given the model's research-only scope, so severity remains high. Because the documentation is complete and detailed, a downstream user will easily see the explanations, yielding a detection score of 9.
SAFETY	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.046, fnr_ratio=0.60. Heuristic metric-to-O/S/D mapping — not validated science.

Reproducibility Information

Model reference	/models/flawed_model_suite/variant2_robustness
Model revision	local
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

No documentation sources recorded for this model.

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_disclosure_gap_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.