

|               |                                                               |
|---------------|---------------------------------------------------------------|
| Model         | /models/flawed_model_suite/variant2_robustness                |
| Evaluation ID | d69972a3-1039-4db9-8578-1227b45657c5                          |
| Status        | FINALIZED                                                     |
| Created       | 2026-10-02T01:34:20.768494Z                                   |
| Report        | report_v1 · version 1 · generated 2026-10-02T17:54:52.905922Z |

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm\_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY\_HEURISTIC\_OSD\_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for /models/flawed\_model\_suite/variant2\_robustness: 6.669/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.669/10 (not FRIES2)
- Human reviewed: no
- Probe flags: empty\_file\_list, evaluation\_contract\_v2, missing\_misuse\_risks, missing\_privacy, missing\_security\_warnings, needs\_human\_review, unpinned\_revision, user\_defined\_local\_dataset
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

6.669 / 10

| Dimension      | Score  |
|----------------|--------|
| SAFETY         | 7.5595 |
| FAIRNESS       | 6.6039 |
| INTEGRITY      | 3.3019 |
| ROBUSTNESS     | 7.5595 |
| EXPLAINABILITY | 8.3203 |

Score type: original\_FRIES · overall confidence 0.9518

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

|                               |                    |
|-------------------------------|--------------------|
| Contract kind                 | documentation_only |
| Task                          | —                  |
| Dataset (comparability label) | —                  |
| Config                        | —                  |
| Pairing ID                    | —                  |

Model Identity

|                 |                                                |
|-----------------|------------------------------------------------|
| Model reference | /models/flawed_model_suite/variant2_robustness |
| Frozen revision | local                                          |

Dataset Identity

|                      |   |
|----------------------|---|
| Dataset key          | — |
| Dataset revision     | — |
| User dataset ID      | — |
| Dataset content hash | — |

Local Execution Environment

|                   |                                          |
|-------------------|------------------------------------------|
| Execution device  | GPU (NVIDIA GeForce RTX 4060 Laptop GPU) |
| Inference backend | local_hf                                 |
| CUDA available    | True                                     |

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

| Aspect         | O (Occurrence) | S (Severity) | D (Detection) |
|----------------|----------------|--------------|---------------|
| FAIRNESS       | 6              | 6            | 8             |
| ROBUSTNESS     | 9              | 6            | 8             |
| INTEGRITY      | 3              | 4            | 3             |
| EXPLAINABILITY | 8              | 8            | 9             |
| SAFETY         | 9              | 6            | 8             |

Source: osd\_agent\_autonomous · human reviewed: no

Confidence summary

Method: geometric\_mean\_v1 · proposed calibration: True

| Scope          | Confidence |
|----------------|------------|
| Overall        | 0.9473     |
| FAIRNESS       | 0.983      |
| ROBUSTNESS     | 1.0        |
| INTEGRITY      | 0.7761     |
| EXPLAINABILITY | 1.0        |
| SAFETY         | 1.0        |

Evidence strength only — not correctness or O/S/D



**FAIRNESS (confidence 0.983)**

evaluation\_contract\_v2 user\_defined\_local\_dataset

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES                                                                                                                                                                                                                                                                                                                                                                                                                   |
| seed                 | 42                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| dp_ci                | {"B": 1000, "ci_lower": 0.48709, "ci_upper": 0.560974, "method": "bootstrap_percentile", "point": 0.526462}                                                                                                                                                                                                                                                                                                                                                                                                                    |
| eo_ci                | {"B": 1000, "ci_lower": 0.302918, "ci_upper": 0.438214, "method": "bootstrap_percentile", "point": 0.370192}                                                                                                                                                                                                                                                                                                                                                                                                                   |
| f1_ci                | {"B": 1000, "ci_lower": 0.24658, "ci_upper": 0.343395, "method": "bootstrap_percentile", "point": 0.294037}                                                                                                                                                                                                                                                                                                                                                                                                                    |
| groups               | {"0": {"f1": 0.526544, "fpr": 0.1875, "n": 2461, "positive_rate": 0.265746, "tpr": 0.903346}, "1": {"f1": 0.82058, "fpr": 0.557692, "n": 539, "positive_rate": 0.792208, "tpr": 0.939577}}                                                                                                                                                                                                                                                                                                                                     |
| inference            | {"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "/models/flawed_model_suite/variant2_robustness", "num_labels": 2, "revision": "local", "task_type": "binary_classification"} |
| model_ref            | /models/flawed_model_suite/variant2_robustness                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| predictor            | inference_backend                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| excess_dpd           | 0.021667                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| dataset_key          | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| min_group_n          | 30                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| n_evaluated          | 3000                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| text_column          | text                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| group_column         | identity_ref                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| missing_rows         | {"missing_count": 0, "missing_reasons": {}}                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| excluded_rows        | {"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| fairness_mode        | user_defined_local                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| label_mapping        | [{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]                                                                                                                                                                                                                                                                                                                                                                                                                               |
| target_column        | label                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| dataset_format       | csv                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| label_encoding       | {"0": 0, "1": 1}                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| label_rate_gap       | 0.504795                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| model_revision       | local                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| excluded_groups      | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| observed_groups      | [{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| dataset_revision     | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| evaluation_class     | user_dataset_v2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| dataset_row_count    | 3000                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| confidence_factors   | {"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}                                                                                                                                                                                                                                                                                                                                                                                                                              |
| dataset_content_id   | c5d8778e-1eba-4a8c-a55c-82487c90a1b2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| inference_executed   | True                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| needs_human_review   | True                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| subgroup_f1_spread   | 0.294037                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| probe_status_reason  | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| sensitive_attribute  | identity_ref                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| min_group_n_observed | 539                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| positive_label_index | 1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |

|                               |          |
|-------------------------------|----------|
| equalized_odds_difference     | 0.370192 |
| demographic_parity_difference | 0.526462 |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| 912e4e59-5b33-4eea-bbfa-1e88208907e3 | sha256:e9715f96dbed052905741e9b2833c335f245e219d90f59516f125eef247262c6 | s3://trustlens/evidence/d69972a3-1039-4db9-8578-1227b45657c5/912e4e59-5b33-4eea-bbfa-1e88208907e3.json |

**ROBUSTNESS (confidence 1.0)**

evaluation\_contract\_v2

user\_defined\_local\_dataset

| Metric              | Value                                                                                                                                                                                                                                                                                                                                         |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| norm                | discrete                                                                                                                                                                                                                                                                                                                                      |
| note                | Layer A evidence only – R-ROB-PERT detection does not assign O/S/D                                                                                                                                                                                                                                                                            |
| seed                | 42                                                                                                                                                                                                                                                                                                                                            |
| attack              | char_swap                                                                                                                                                                                                                                                                                                                                     |
| status              | EVALUATED                                                                                                                                                                                                                                                                                                                                     |
| dataset             | {"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}                                                                                                                                                                                                                                                      |
| epsilon             | 0.03                                                                                                                                                                                                                                                                                                                                          |
| inference           | {"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}      |
| model_ref           | /models/flawed_model_suite/variant2_robustness                                                                                                                                                                                                                                                                                                |
| n_samples           | 3000                                                                                                                                                                                                                                                                                                                                          |
| task_type           | null                                                                                                                                                                                                                                                                                                                                          |
| dataset_key         | null                                                                                                                                                                                                                                                                                                                                          |
| limitations         | []                                                                                                                                                                                                                                                                                                                                            |
| max_changes         | 3                                                                                                                                                                                                                                                                                                                                             |
| n_evaluated         | 3000                                                                                                                                                                                                                                                                                                                                          |
| n_requested         | 3000                                                                                                                                                                                                                                                                                                                                          |
| reliability         | {"failed_gates": [], "gates_passed": true}                                                                                                                                                                                                                                                                                                    |
| text_column         | text                                                                                                                                                                                                                                                                                                                                          |
| uncertainty         | {"accuracy_drop": {"B": 1000, "ci_lower": 0.03, "ci_upper": 0.052667, "method": "bootstrap_percentile", "point": 0.041}, "clean_accuracy": {"ci_lower": 0.794542, "ci_upper": 0.822668, "method": "wilson_score", "point": 0.809}, "robust_accuracy": {"ci_lower": 0.752558, "ci_upper": 0.782756, "method": "wilson_score", "point": 0.768}} |
| probe_status        | EVALUATED                                                                                                                                                                                                                                                                                                                                     |
| accuracy_drop       | 0.041                                                                                                                                                                                                                                                                                                                                         |
| excluded_rows       | {"missing_target": 0, "missing_text": 0, "unmapped_label": 0}                                                                                                                                                                                                                                                                                 |
| label_mapping       | [{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]                                                                                                                                                                                                                                              |
| target_column       | label                                                                                                                                                                                                                                                                                                                                         |
| aspect_scoring      | no_material_risk                                                                                                                                                                                                                                                                                                                              |
| clean_accuracy      | 0.809                                                                                                                                                                                                                                                                                                                                         |
| dataset_format      | csv                                                                                                                                                                                                                                                                                                                                           |
| label_encoding      | {"0": 0, "1": 1}                                                                                                                                                                                                                                                                                                                              |
| model_revision      | local                                                                                                                                                                                                                                                                                                                                         |
| scored_risk_id      | null                                                                                                                                                                                                                                                                                                                                          |
| risks_triggered     | []                                                                                                                                                                                                                                                                                                                                            |
| robust_accuracy     | 0.768                                                                                                                                                                                                                                                                                                                                         |
| dataset_revision    | null                                                                                                                                                                                                                                                                                                                                          |
| evaluation_class    | user_dataset_v2                                                                                                                                                                                                                                                                                                                               |
| n_perturb_failed    | 0                                                                                                                                                                                                                                                                                                                                             |
| proposed_mapping    | False                                                                                                                                                                                                                                                                                                                                         |
| degradation_ratio   | 0.9493                                                                                                                                                                                                                                                                                                                                        |
| confidence_factors  | {"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}                                                                                                                                                                                                                                                |
| dataset_content_id  | c5d8778e-1eba-4a8c-a55c-82487c90a1b2                                                                                                                                                                                                                                                                                                          |
| inference_executed  | True                                                                                                                                                                                                                                                                                                                                          |
| n_label_compatible  | 3000                                                                                                                                                                                                                                                                                                                                          |
| attack_success_rate | 0.0733                                                                                                                                                                                                                                                                                                                                        |

|                          |                                                                                       |
|--------------------------|---------------------------------------------------------------------------------------|
| methodology_version      | tl-methodology-v1.0                                                                   |
| probe_status_reason      | no material perturbation risk at 95% CI (drop=0.041, ci_lower=0.03, $\epsilon$ =0.05) |
| relative_degradation     | 0.9493                                                                                |
| label_compat_fraction    | 1.0                                                                                   |
| perturbation_coverage    | 1.0                                                                                   |
| n_successfully_perturbed | 3000                                                                                  |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                 |
|--------------------------------------|-------------------------------------------------------------------------|-----------------------------------------------------|
| 50784310-8830-4809-9ce3-824595c3d4f4 | sha256:a21e687d0dfbe202c86074b5ffaaf30913edc3059b3014b3991a4b70a7da5919 | s3://trustlens/evidence/<br>d69972a3-1039-4db9-8578 |

INTEGRITY (confidence 0.7761)

unpinned\_revision

empty\_file\_list

| Metric              | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|---------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                | Layer A evidence only – Integrity risks do not assign O/S/D                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| checks              | <pre>{"card_present": {"detail": "1809 chars", "pass": true}, "files_listed": {"detail": "empty or missing files list", "pass": false}, "files_listing_recorded": {"detail": "empty file list \u2014 no listing fingerprint identity", "files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "model_revision": "local", "pass": false}, "license_declared": {"detail": "apache-2.0", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training_data, seed, evaluation, hyperparameters", "pass": true}, "revision_pinned": {"detail": "non-sha-like revision=local", "pass": false}}</pre>                                                                                                                  |
| claims              | <pre>{"cryptographic_identity": {"reason": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "files_listing_currency": {"reason": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "identity_recording": {"reason": "Hub file manifest missing \u2014 listing-dependent identity claims insufficient", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre> |
| status              | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| identity            | <pre>{"files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "files_listing_reverification": "not performed", "hash_comparison": "not performed", "hub_files": [], "live_files_fingerprint": null, "revision": "local", "sha_like": false, "trusted_reference": null, "weight_hash": null}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                |
| disclosure          | <pre>{"card_chars": 1809, "card_present": true, "license_structured": "apache-2.0", "repro_groups_found": ["training_data", "seed", "evaluation", "hyperparameters"]}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| limitations         | <p>[ "Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files." ]</p>    |
| reliability         | <pre>{"failed_gates": ["G-INT-LISTING-UNVERIFIED", "G-INT-HASH-REF-MISSING", "G-INT-HASH-LOCAL-MISSING", "I-INT-HASH-UNVERIFIED"], "gates_passed": false}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| uncertainty         | {}                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| probe_status        | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| osd_proposals       | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| aspect_scoring      | risk_detected                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| claim_boundary      | <pre>{"cryptographic_identity": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "files_listing_currency": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim."}</pre>                                                                                                                                                                                                                                                                                                                                  |
| scored_risk_id      | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| risks_triggered     | ["I-INT-REV-UNPINNED", "I-INT-MANIFEST-MISSING"]                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| proposed_mapping    | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| methodology_basis   | TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| confidence_factors  | <pre>{"combined": 0.7761, "data_quality": 0.55, "evidence_completeness": 0.85, "probe_reliability": 1.0}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| methodology_version | tl-integrity-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| ef1f1be2-84f8-4b04-a0c2-0733f04e29a8 | sha256:fe789b49166a7672c5a1382a893dc4971d059ee4ac540b18560c857b6f1364b6 | s3://trustlens/evidence/d69972a3-1039-4db9-8578-1227b45657c5/ef1f1be2-84f8-4b04-a0c2-0733f04e29a8.json |

**EXPLAINABILITY (confidence 1.0)**

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|----------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Layer A documentation/transparency evidence only – not O/S/D or interpretability quality                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| checks               | <pre>{"documentation_completeness": {"coverage_ratio": 1.0, "detail": "5/5 required sections present", "pass": true, "sections_present": 5, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "Ethical Considerations", "pass": true}, "evaluation": {"detail": "Evaluation", "pass": true}, "intended_use": {"detail": "Intended Use", "pass": true}, "limitations": {"detail": "Limitations", "pass": true}, "pass": true, "training_data": {"detail": "Training Data", "pass": true}}</pre> |
| claims               | <pre>{"documentation_completeness": {"reason": "documentation checklist evaluated; 5/5 required sections present", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</pre>                                                                                                                                                                                                                                                                                                                                                       |
| status               | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| sections             | <pre>{"ethical_considerations": {"matched_heading": "Ethical Considerations", "present": true}, "evaluation": {"matched_heading": "Evaluation", "present": true}, "intended_use": {"matched_heading": "Intended Use", "present": true}, "limitations": {"matched_heading": "Limitations", "present": true}, "training_data": {"matched_heading": "Training Data", "present": true}}</pre>                                                                                                                                                                                                                   |
| card_chars           | 1809                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| limitations          | ["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]                          |
| reliability          | <pre>{"failed_gates": [], "gates_passed": true}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| osd_proposals        | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| aspect_scoring       | no_material_risk                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| bonus_sections       | <pre>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                  |
| claim_boundary       | <pre>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</pre>                                                                                                                                                          |
| contradictions       | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| coverage_ratio       | 1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| scored_risk_id       | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| risks_triggered      | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| sections_present     | 5                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| methodology_basis    | TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| sections_required    | 5                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| confidence_factors   | <pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| methodology_version  | tl-explainability-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| documentation_source | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                             |
|--------------------------------------|-------------------------------------------------------------------------|-------------------------------------------------|
| 50da83ab-0f80-4554-98e2-73f7a9c59a16 | sha256:6d13cde42890974cbd84530bf183b0c162a2fd329d9a3addf85072dce5da8292 | s3://trustlens/evidence/d69972a3-1039-4db9-8578 |

**SAFETY (confidence 1.0)**

missing\_misuse\_risks

missing\_privacy

missing\_security\_warnings

needs\_human\_review

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| checks               | <pre>{"data_disclosure": {"detail": "Training Data", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</pre>                                                                                                                                                                                                                                                                      |
| claims               | <pre>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| status               | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| behavior             | <pre>{"benign_fpr": 0.2195833333333332, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.6012950971322849, "harmful_recall": 0.9233333333333333, "min_severe_n": 30, "model_ref": "/models/flawed_model_suite/variant2_robustness", "model_revision": "local", "n_evaluated": 3000, "overall_fnr": 0.07666666666666666, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.04609929078014184, "severe_fnr_ci": [0.02713545260229284, 0.07726362716776161], "severe_n": 282, "status": "EVALUATED"}</pre>                                                           |
| fnr_ratio            | 0.6012950971322849                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| card_chars           | 1809                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| severe_fnr           | 0.04609929078014184                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| limitations          | <p>[ "Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned.", "aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score." ]</p> |
| reliability          | <pre>{"failed_gates": [], "gates_passed": true}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| uncertainty          | <pre>{}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| bonus_checks         | <pre>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| osd_proposals        | <pre>[]</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| aspect_scoring       | risk_detected                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| checks_present       | 1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| claim_boundary       | <pre>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</pre>                                                                       |
| coverage_ratio       | 0.25                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| scored_risk_id       | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| checks_required      | 4                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| required_checks      | <pre>{"data_disclosure": {"matched_heading": "Training Data", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| risks_triggered      | <pre>[ "S-GOV-DISCLOSURE-GAP" ]</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| methodology_basis    | TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| confidence_factors   | <pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| high_impact_claims   | <pre>[]</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| methodology_version  | tl-safety-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| documentation_source | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| d3be59fd-0015-4a7f-b175-fd36a1784d58 | sha256:da6604eded297da1e26cce23310754405ebd2ed2eeff88153bd88245b49563b8 | s3://trustlens/evidence/d69972a3-1039-4db9-8578-1227b45657c5/d3be59fd-0015-4a7f-b175-fd36a1784d58.json |

## Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

### FAIRNESS

|                       |                           |
|-----------------------|---------------------------|
| Status                | EVALUATED                 |
| Risk decision         | Not scored                |
| Gate / rule           | No gate failures recorded |
| Human O/S/D           | O=6 · S=6 · D=8           |
| FRIES dimension score | 6.604                     |

### ROBUSTNESS

|                       |                                                                                         |
|-----------------------|-----------------------------------------------------------------------------------------|
| Status                | EVALUATED — no material perturbation risk at 95% CI (drop=0.041, ci_lower=0.03, ε=0.05) |
| Risk decision         | no material risk                                                                        |
| Gate / rule           | No gate failures recorded                                                               |
| Human O/S/D           | O=9 · S=6 · D=8                                                                         |
| FRIES dimension score | 7.559                                                                                   |

### INTEGRITY

|                       |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|-----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Status                | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Risk decision         | risk detected — risks: I-INT-REV-UNPINNED, I-INT-MANIFEST-MISSING                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| Gate / rule           | G-INT-LISTING-UNVERIFIED, G-INT-HASH-REF-MISSING, G-INT-HASH-LOCAL-MISSING, I-INT-HASH-UNVERIFIED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| Human O/S/D           | O=3 · S=4 · D=3                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| FRIES dimension score | 3.302                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| Limitations           | Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files. |

### EXPLAINABILITY

|                        |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Status                 | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Risk decision          | no material risk                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| Gate / rule            | No gate failures recorded                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Documentation coverage | 1.00                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| Human O/S/D            | O=8 · S=8 · D=9                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| FRIES dimension score  | 8.320                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| Limitations            | Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score. |

## SAFETY

|                        |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Status                 | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Risk decision          | risk detected — risks: S-GOV-DISCLOSURE-GAP                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| Gate / rule            | No gate failures recorded                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Documentation coverage | 0.25                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Human O/S/D            | O=9 · S=6 · D=8                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| FRIES dimension score  | 7.559                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| Limitations            | Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score. |

## O/S/D representation

Methodology status: **LEGACY\_HEURISTIC\_OSD\_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

| Aspect         | O<br>(Occurrence) | S<br>(Severity) | D<br>(Detection) | Evidence<br>strength | Rationale                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|----------------|-------------------|-----------------|------------------|----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| FAIRNESS       | 6                 | 6               | 8                | 0.983                | [PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.370. Heuristic metric-to-O/S/D mapping — not validated science.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| ROBUSTNESS     | 9                 | 6               | 8                | 1.0                  | [PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.809, robust_accuracy=0.768, accuracy_drop=0.041. Heuristic metric-to-O/S/D mapping — not validated science.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| INTEGRITY      | 3                 | 4               | 3                | 0.7761               | The integrity probe flags two concrete issues: an "unpinned_revision" (revision reported as "local" and not SHA-like) and an "empty_file_list" (no files listed, fingerprint is the hash of an empty string). Both flags indicate that the model's identity cannot be reliably reproduced, which is a clear integrity risk. License disclosure is present, but the missing cryptographic identity and unverified file-listing mean the risk is not mitigated. Consequently the occurrence of the risk is low (O=3), the potential consequence of using an unverified model is moderate (S=4), and a downstream user would likely notice the missing revision and file list only after inspecting the hub metadata (D=3).                |
| EXPLAINABILITY | 8                 | 8               | 9                | 1.0                  | All five required documentation sections are present and contain model-specific details: the Intended Use explains the classifier's offline benchmarking purpose; Limitations describe the deliberately injected robustness flaw and overfitting; Training Data specifies the 600-row sample and seed; Evaluation reports a 3,000-row test set and a character-swap probe; Ethical Considerations warn against real-world moderation use. The text is concrete and directly tied to this model rather than generic boiler-plate, satisfying substantive explainability. Because the risk of missing explainability is essentially absent, occurrence is high (O=8) and both severity and detection are correspondingly high (S=8, D=9). |
| SAFETY         | 9                 | 6               | 8                | 1.0                  | [PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.046, fnr_ratio=0.60. Heuristic metric-to-O/S/D mapping — not validated science.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |

OSD overall evidence-strength mean: 0.9518 (not FRIES certainty)

## Reproducibility Information

|                   |                                                |
|-------------------|------------------------------------------------|
| Model reference   | /models/flawed_model_suite/variant2_robustness |
| Model revision    | local                                          |
| TrustLens version | —                                              |
| Contract kind     | documentation_only                             |

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

## Documentation Sources

No documentation sources recorded for this model.

## Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-\* values are evidence-layer detections only; scored\_risk\_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored\_risk\_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect\_scoring=risk\_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect\_scoring=governance\_disclosure\_gap means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage\_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage\_ratio counts present required sections; it is not an explainability quality score.
- high\_impact\_claims are documentation metadata flags; they are never promoted to risks\_triggered and do not map to O/S/D or FRIES.

report\_v1 · report version 1 · generated 2026-10-02T17:54:52.905922Z · canonical JSON is the source of truth; this PDF is a rendered projection.