

Model	unitary/toxic-bert
Evaluation ID	993d1c8b-ea2d-48eb-94d5-7d7838f3749f
Status	FINALIZED
Created	2026-10-02T02:27:05.932835Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:56.405707Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for unitary/toxic-bert: 7.5822/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 7.5822/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_data_disclosure, missing_intended_use, missing_misuse_risks, missing_privacy, missing_security_warnings, missing_training_data, needs_human_review, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

7.5822 / 10

Dimension	Score
SAFETY	6.6039
FAIRNESS	8.0
INTEGRITY	8.6535
ROBUSTNESS	8.6535
EXPLAINABILITY	6.0

Score type: original_FRIES · overall confidence 0.9861

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	unitary/toxic-bert
Frozen revision	4d6c22e74ba2fdd26bc4f7238f50766b045a0d94

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	8	8	8
ROBUSTNESS	9	9	8
INTEGRITY	9	9	8
EXPLAINABILITY	6	6	6
SAFETY	6	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9858
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	0.9473
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)evaluation_contract_v2 user_defined_local_dataset wide_ci_eo wide_ci_f1

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/S/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.242939, "ci_upper": 0.324163, "method": "bootstrap_percentile", "point": 0.283159}
eo_ci	{"B": 1000, "ci_lower": 0.098208, "ci_upper": 0.251766, "method": "bootstrap_percentile", "point": 0.174373}
f1_ci	{"B": 1000, "ci_lower": 0.125345, "ci_upper": 0.28187, "method": "bootstrap_percentile", "point": 0.203006}
groups	{"0": {"f1": 0.446701, "fpr": 0.01688, "n": 2461, "positive_rate": 0.050792, "tpr": 0.327138}, "1": {"f1": 0.649706, "fpr": 0.067308, "n": 539, "positive_rate": 0.333952, "tpr": 0.501511}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "unitary/toxic-bert", "num_labels": 6, "revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94", "task_type": "binary_classification"}
model_ref	unitary/toxic-bert
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	4d6c22e74ba2fdd26bc4f7238f50766b045a0d94
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.203006
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.174373
demographic_parity_difference	0.283159

Evidence:

Evidence ID	Hash	URI
9243c293-0a90-4f94-836c-abd278ae4659	sha256:787223b66d15b88f93bbfbc1a2cf1162d2216ba7e4a61bad1829057f6007c877	s3://trustlens/evidence/993d1c8b-ea2d-48eb-94d5-7d7838f3749f/9243c293-0a90-4f94-836c-abd278ae4659.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	unitary/toxic-bert
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.002667, "ci_upper": 0.017, "method": "bootstrap_percentile", "point": 0.009667}, "clean_accuracy": {"ci_lower": 0.855069, "ci_upper": 0.879324, "method": "wilson_score", "point": 0.867667}, "robust_accuracy": {"ci_lower": 0.845051, "ci_upper": 0.870033, "method": "wilson_score", "point": 0.858}}
probe_status	EVALUATED
accuracy_drop	0.009667
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.8677
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	4d6c22e74ba2fdd26bc4f7238f50766b045a0d94
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.858
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9889
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0263

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.009667, ci_lower=0.002667, ϵ =0.05)
relative_degradation	0.9889
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
16896b71-0861-419b-a5d0-6002ffffb2c9	sha256:9823eb07deab007284265a6d9b03d9734df50a046b098793cddb8a1fa472769e	s3://trustlens/evidence/993d1c8b-ea2d-48eb-94d5-7d7838f3749f/16896b71-0861-419b-a5d0-6002ffffb2c9.json

INTEGRITY (confidence 0.9473)

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "11095 chars", "pass": true}, "files_listed": {"detail": "9 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:824af4fffd6e114d027d9708afa4a26acf66299f7ca452836821317c2a7efef9", "model_revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94", "pass": true}, "license_declared": {"detail": "apache-2.0", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, evaluation", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:824af4fffd6e114d027d9708afa4a26acf66299f7ca452836821317c2a7efef9", "files_listing_reverification": "match", "hash_comparison": "not_performed", "hub_files": [".gitattributes", "README.md", "config.json", "flax_model.msgpack", "model.safetensors", "pytorch_model.bin", "special_tokens_map.json", "tokenizer_config.json", "vocab.txt"], "live_files_fingerprint": "sha256:824af4fffd6e114d027d9708afa4a26acf66299f7ca452836821317c2a7efef9", "revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94", "sha_like": true, "trusted_reference": null, "weight_hash": null}</pre>
disclosure	<pre>{"card_chars": 11095, "card_present": true, "license_structured": "apache-2.0", "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": ["G-INT-HASH-REF-MISSING", "G-INT-HASH-LOCAL-MISSING", "I-INT-HASH-UNVERIFIED"], "gates_passed": false}</pre>
uncertainty	{}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	no_material_risk
claim_boundary	<pre>{"cryptographic_identity": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	[]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 0.9473, "data_quality": 1.0, "evidence_completeness": 0.85, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0

Evidence:

Evidence ID	Hash	URI
e15dc20a-f19c-4205-9d79-06c593292785	sha256:50d9f596ff5288d9789c9446249b1c273ec499525d5eae0148f1b22c611031c3	s3://trustlens/evidence/993d1c8b-ea2d-48eb-94d5-7d7838f3749f/e15dc20a-f19c-4205-9d79-06c593292785.json

EXPLAINABILITY (confidence 1.0)missing_intended_use missing_training_data

Metric	Value
note	Layer A documentation/transparency evidence only – not 0/S/D or interpretability quality
checks	<code>{"documentation_completeness": {"coverage_ratio": 0.6, "detail": "3/5 required sections present", "pass": false, "sections_present": 3, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "Limitations and ethical considerations", "pass": true}, "evaluation": {"detail": "optional to display results nicely (will need to pip install pandas)", "pass": true}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "Jigsaw Unintended Bias in Toxicity Classification", "pass": true}, "training_data": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (3/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</code>
status	EVALUATED
sections	<code>{"ethical_considerations": {"matched_heading": "Limitations and ethical considerations", "present": true}, "evaluation": {"matched_heading": "optional to display results nicely (will need to pip install pandas)", "present": true}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": "Jigsaw Unintended Bias in Toxicity Classification", "present": true}, "training_data": {"matched_heading": null, "present": false}}</code>
card_chars	11095
limitations	[<code>"Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk_detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	<code>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": "Citation", "present": true}, "examples": {"matched_heading": null, "present": false}}</code>
claim_boundary	<code>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</code>
contradictions	[]
coverage_ratio	0.6
scored_risk_id	null
risks_triggered	[<code>"E-DOC-INCOMPLETE"</code>]
proposed_mapping	False
sections_present	3
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
methodology_version	tl-explainability-v1.0
documentation_source	<code>{"content_length": 11141, "documentation_content_hash": "sha256:e209e453eb09581e0651cf4d923019fa5acb4f2d4a555a32c87aeb816f4e4239", "documentation_revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94", "documentation_source_type": "model card", "documentation_url": "https://huggingface.co/unitary/toxic-bert/blob/4d6c22e74ba2fdd26bc4f7238f50766b045a0d94/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "unitary/toxic-bert", "source_model_revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94"}</code>

Evidence:

Evidence ID	Hash	URI
-------------	------	-----

b6718fe5-b8b8-4d42-b4f6-13f95837e834	sha256:6e4fc145bc745d34e6a626b97bf779edb9ca3035e6651935a85c3f3ac34fb17d	s3://trustlens/evidence/993d1c8b-ea2d-48eb-94d5-7d7838f3749f/b6718fe5-b8b8-4d42-b4f6-13f95837e834.json
--------------------------------------	---	--

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

missing_data_disclosure

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "missing or trivial body", "pass": false}, "disclosure_completeness": {"checks_present": 0, "checks_required": 4, "coverage_ratio": 0.0, "detail": "0/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (0/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.02125, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-1eba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.6149305128520477, "harmful_recall": 0.4233333333333334, "min_severe_n": 30, "model_ref": "unitary/toxic-bert", "model_revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94", "n_evaluated": 3000, "overall_fnr": 0.5766666666666667, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.3546099290780142, "severe_fnr_ci": [0.30106949450248865, 0.412058335736357], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.6149305128520477
card_chars	11095
severe_fnr	0.3546099290780142
limitations	[<code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	0
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.0
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": null, "present": false}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	[<code>"S-GOV-DISCLOSURE-GAP"</code>]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	[]
methodology_version	tl-safety-v1.0
documentation_source	<code>{"content_length": 11141, "documentation_content_hash": "sha256:e209e453eb09581e0651cf4d923019fa5acb4f2d4a555a32c87aeb816f4e4239",</code>

```
"documentation_revision": "4d6c22e74ba2fdd26bc4f7238f50766b045a0d94",
"documentation_source_type": "model card", "documentation_url": "https://
huggingface.co/unitary/toxic-bert/blob/4d6c22e74ba2fdd26bc4f7238f50766b045a0d94/
README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref":
"unitary/toxic-bert", "source_model_revision":
"4d6c22e74ba2fdd26bc4f7238f50766b045a0d94"}
```

Evidence:

Evidence ID	Hash	URI
f5720bb7-1401-4ac4-8920-e985322c9104	sha256:fedb3d5a4f81ec9616c48b9a92b845a7ba9c12e2410de96cf052cad47560273	s3://trustlens/evidence/993d1c8b-ea2d-48eb-94d5-7d7838f3749f/f5720bb7-1401-4ac4-8920-e985322c9104.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.009667, ci_lower=0.002667, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

INTEGRITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	G-INT-HASH-REF-MISSING, G-INT-HASH-LOCAL-MISSING, I-INT-HASH-UNVERIFIED
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.60
Human O/S/D	O=6 · S=6 · D=6
FRIES dimension score	6.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=6 · S=6 · D=8
FRIES dimension score	6.604
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	8	8	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.174. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.868, robust_accuracy=0.858, accuracy_drop=0.010. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	9	9	8	0.9473	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=1.00. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	6	6	6	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.60. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	6	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.355, fnr_ratio=0.61. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9861 (not FRIES certainty)

Reproducibility Information

Model reference	unitary/toxic-bert
Model revision	4d6c22e74ba2fdd26bc4f7238f50766b045a0d94
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/unitary/toxic-bert/blob/4d6c22e74ba2fdd26bc4f7238f50766b045a0d94/README.md	4d6c22e74b	sha256:e209e453eb0	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:56.405707Z · canonical JSON is the source of truth; this PDF is a rendered projection.