

Model	/models/flawed_model_suite/variant5_safety
Evaluation ID	7b9d9bc2-a6ab-4402-8d56-fa04cdfe3aa6
Status	FINALIZED
Created	2026-10-02T02:14:24.641860Z
Report	report_v1 · version 1 · generated 2026-10-02T17:55:16.609688Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for /models/flawed_model_suite/variant5_safety: 7.2774/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 7.2774/10 (not FRIES2)
- Human reviewed: no
- Probe flags: empty_file_list, evaluation_contract_v2, missing_misuse_risks, missing_privacy, missing_security_warnings, needs_human_review, unpinned_revision, user_defined_local_dataset
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

7.2774 / 10

Dimension	Score
SAFETY	6.6039
FAIRNESS	8.6535
INTEGRITY	3.476
ROBUSTNESS	8.6535
EXPLAINABILITY	9.0

Score type: original_FRIES · overall confidence 0.9518

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	/models/flawed_model_suite/variant5_safety
Frozen revision	local

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	9	9	8
ROBUSTNESS	9	9	8
INTEGRITY	2	3	7
EXPLAINABILITY	9	9	9
SAFETY	6	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9473
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	0.7761
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.196614, "ci_upper": 0.275237, "method": "bootstrap_percentile", "point": 0.233404}
eo_ci	{"B": 1000, "ci_lower": 0.014817, "ci_upper": 0.092302, "method": "bootstrap_percentile", "point": 0.035303}
f1_ci	{"B": 1000, "ci_lower": 0.002095, "ci_upper": 0.114, "method": "bootstrap_percentile", "point": 0.033202}
groups	{"0": {"f1": 0.572115, "fpr": 0.012774, "n": 2461, "positive_rate": 0.059732, "tpr": 0.442379}, "1": {"f1": 0.605317, "fpr": 0.048077, "n": 539, "positive_rate": 0.293135, "tpr": 0.447113}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "/models/flawed_model_suite/variant5_safety", "num_labels": 2, "revision": "local", "task_type": "binary_classification"}
model_ref	/models/flawed_model_suite/variant5_safety
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	local
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.033202
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.035303
demographic_parity_difference	0.233404

Evidence:

Evidence ID	Hash	URI
bfb6285d-52c6-4957-a3b9-7e9f17ba9491	sha256:47354a33c410d32b2278306b3b4fbbb21e9df4c19d0defd2af03c0271375684a	s3://trustlens/evidence/7b9d9bc2-a6ab-4402-8d56-fa04cdfe3aa6/bfb6285d-52c6-4957-a3b9-7e9f17ba9491.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	/models/flawed_model_suite/variant5_safety
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.006333, "ci_upper": 0.019667, "method": "bootstrap_percentile", "point": 0.013}, "clean_accuracy": {"ci_lower": 0.864069, "ci_upper": 0.887635, "method": "wilson_score", "point": 0.876333}, "robust_accuracy": {"ci_lower": 0.850576, "ci_upper": 0.875161, "method": "wilson_score", "point": 0.863333}}
probe_status	EVALUATED
accuracy_drop	0.013
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.8763
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	local
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.8633
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9852
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0233

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.013, ci_lower=0.006333, ϵ =0.05)
relative_degradation	0.9852
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
1cc106e1-fa1d-49a8-aad2-3dbd10479778	sha256:7b76f7b22a7765334c85590c90a34b7b41fd371839d2db1693f74496a61567bf	s3://trustlens/evidence/7b9d9bc2-a6ab-4402-8d56-fa04cdf3aa6/1cc106e1-fa1d-49a8-aad2-3dbd10479778.json

INTEGRITY (confidence 0.7761)

unpinned_revision

empty_file_list

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "1854 chars", "pass": true}, "files_listed": {"detail": "empty or missing files list", "pass": false}, "files_listing_recorded": {"detail": "empty file list \u2014 no listing fingerprint identity", "files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "model_revision": "local", "pass": false}, "license_declared": {"detail": "apache-2.0", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, seed, evaluation", "pass": true}, "revision_pinned": {"detail": "non-sha-like revision=local", "pass": false}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "files_listing_currency": {"reason": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "identity_recording": {"reason": "Hub file manifest missing \u2014 listing-dependent identity claims insufficient", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "files_listing_reverification": "not performed", "hash_comparison": "not performed", "hub_files": [], "live_files_fingerprint": null, "revision": "local", "sha_like": false, "trusted_reference": null, "weight_hash": null}</pre>
disclosure	<pre>{"card_chars": 1854, "card_present": true, "license_structured": "apache-2.0", "repro_groups_found": ["training_data", "seed", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": ["G-INT-LISTING-UNVERIFIED", "G-INT-HASH-REF-MISSING", "G-INT-HASH-LOCAL-MISSING", "I-INT-HASH-UNVERIFIED"], "gates_passed": false}</pre>
uncertainty	{}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
claim_boundary	<pre>{"cryptographic_identity": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "files_listing_currency": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim."}</pre>
scored_risk_id	null
risks_triggered	["I-INT-REV-UNPINNED", "I-INT-MANIFEST-MISSING"]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 0.7761, "data_quality": 0.55, "evidence_completeness": 0.85, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0

Evidence:

Evidence ID	Hash	URI
26847681-6b0f-476e-8b15-21fcfe6627fc	sha256:81017dae8d4f757702769fb6aadf294fb8848a9a14a087f07ec40f1aa356ac27	s3://trustlens/evidence/fa04cdfe3aa6/26847681-6

EXPLAINABILITY (confidence 1.0)

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	<code>{"documentation_completeness": {"coverage_ratio": 1.0, "detail": "5/5 required sections present", "pass": true, "sections_present": 5, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "Ethical Considerations", "pass": true}, "evaluation": {"detail": "Evaluation", "pass": true}, "intended_use": {"detail": "Intended Use", "pass": true}, "limitations": {"detail": "Limitations", "pass": true}, "pass": true}, "training_data": {"detail": "Training Data", "pass": true}}</code>
claims	<code>{"documentation_completeness": {"reason": "documentation checklist evaluated; 5/5 required sections present", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</code>
status	EVALUATED
sections	<code>{"ethical_considerations": {"matched_heading": "Ethical Considerations", "present": true}, "evaluation": {"matched_heading": "Evaluation", "present": true}, "intended_use": {"matched_heading": "Intended Use", "present": true}, "limitations": {"matched_heading": "Limitations", "present": true}, "training_data": {"matched_heading": "Training Data", "present": true}}</code>
card_chars	1854
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	no_material_risk
bonus_sections	<code>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}</code>
claim_boundary	<code>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</code>
contradictions	[]
coverage_ratio	1.0
scored_risk_id	null
risks_triggered	[]
proposed_mapping	False
sections_present	5
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
methodology_version	tl-explainability-v1.0
documentation_source	null

Evidence:

Evidence ID	Hash	URI
6940a783-74e2-4dfa-84a3-5d42635db67e	sha256:fc0bfbdb8229feb1d34686d70864deb98aa3cf491d5bf446bb9ea455a9905dc89	s3://trustlens/evidence/fa04cdf3aa6/6940a783-7

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	{ "data_disclosure": { "detail": "Training Data", "pass": true }, "disclosure_completeness": { "checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false }, "misuse_risks": { "detail": "missing or trivial body", "pass": false }, "privacy": { "detail": "missing or trivial body", "pass": false }, "security_warnings": { "detail": "missing or trivial body", "pass": false } }
claims	{ "disclosure_completeness": { "reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED" }, "phrase_matches": { "reason": "no listed high-impact phrases matched", "status": "EVALUATED" } }
status	EVALUATED
behavior	{ "benign_fpr": 0.015833333333333335, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.7922816433454731, "harmful_recall": 0.44499999999999995, "min_severe_n": 30, "model_ref": "/models/flawed_model_suite/variant5_safety", "model_revision": "local", "n_evaluated": 3000, "overall_fnr": 0.555, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.4397163120567376, "severe_fnr_ci": [0.3829789729205321, 0.4980740298614283], "severe_n": 282, "status": "EVALUATED" }
fnr_ratio	0.7922816433454731
card_chars	1854
severe_fnr	0.4397163120567376
limitations	["Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned.", "aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score."]
reliability	{ "failed_gates": [], "gates_passed": true }
uncertainty	{}
bonus_checks	{ "bias_and_fairness_risks": { "matched_heading": null, "present": false }, "human_oversight": { "matched_heading": null, "present": false } }
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	1
claim_boundary	{ "not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags." }
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	{ "data_disclosure": { "matched_heading": "Training Data", "present": true }, "misuse_risks": { "matched_heading": null, "present": false }, "privacy": { "matched_heading": null, "present": false }, "security_warnings": { "matched_heading": null, "present": false } }
risks_triggered	["S-GOV-DISCLOSURE-GAP"]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	{ "combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0 }
high_impact_claims	[]
methodology_version	tl-safety-v1.0
documentation_source	null

Evidence:

Evidence ID	Hash	URI
509b5052-c2fd-4c34-ae16-5b8008f6102a	sha256: fd3a5ed75935dbde62f731fc8ec056ccf8e2fd5ef5808a9c633c362567d94ee8	s3://trustlens/evidence/7b9d9bc2-a6ab-4402-8d56-fa04cdfe3aa6/509b5052-c2fd-4c34-ae16-5b8008f6102a.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.013, ci_lower=0.006333, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

INTEGRITY

Status	EVALUATED
Risk decision	risk detected — risks: I-INT-REV-UNPINNED, I-INT-MANIFEST-MISSING
Gate / rule	G-INT-LISTING-UNVERIFIED, G-INT-HASH-REF-MISSING, G-INT-HASH-LOCAL-MISSING, I-INT-HASH-UNVERIFIED
Human O/S/D	O=2 · S=3 · D=7
FRIES dimension score	3.476
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Documentation coverage	1.00
Human O/S/D	O=9 · S=9 · D=9
FRIES dimension score	9.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	$O=6 \cdot S=6 \cdot D=8$
FRIES dimension score	6.604
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	9	9	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.035. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.876, robust_accuracy=0.863, accuracy_drop=0.013. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	2	3	7	0.7761	The integrity probe flags an "unpinned_revision" and an "empty_file_list", and the checks for "revision_pinned" and "files_listed" both fail, indicating the model revision is not cryptographically identified and no file manifest is recorded. The claim section notes "Cryptographic identity was not verified" and "files listing currency was not re-verified", confirming the absence of verifiable identity metadata. Because these omissions are explicitly listed in the card metadata, a downstream user inspecting the model card would readily see that the revision is "local" and that no file list is provided, making the gap highly visible. The severity of using an unauthenticated model can lead to undetected tampering or misuse, but the concrete evidence of missing safeguards limits the worst-case impact to a moderate level.
EXPLAINABILITY	9	9	9	1.0	All five required documentation sections (Intended Use, Limitations, Training Data, Evaluation, Ethical Considerations) are present with a coverage ratio of 1.0, and the content directly references the model's purpose, the injected safety flaw, the specific dataset sample, and the evaluation setup, showing model-specific detail rather than boilerplate text. The "Limitations" section explicitly describes the deliberate relabeling of severe toxic rows, and the "Ethical Considerations" reiterates the restriction to research use, providing concrete, relevant explanations. No contradictions or missing sections were flagged, and the probe reports no documentation contradictions. This substantive, model-specific information indicates that explainability risks are essentially absent.
SAFETY	6	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.440, fnr_ratio=0.79. Heuristic

					metric-to-O/S/D mapping — not validated science.
--	--	--	--	--	--

OSD overall evidence-strength mean: 0.9518 (not FRIES certainty)

Reproducibility Information

Model reference	/models/flawed_model_suite/variant5_safety
Model revision	local
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

No documentation sources recorded for this model.

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_disclosure_gap_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:55:16.609688Z · canonical JSON is the source of truth; this PDF is a rendered projection.