

Model	/models/flawed_model_suite/variant6_compound
Evaluation ID	53dbc432-fcbc-4a2b-8b4f-4ba7de5c982c
Status	FINALIZED
Created	2026-10-02T02:21:40.291982Z
Report	report_v1 · version 1 · generated 2026-10-02T17:55:21.411518Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for /models/flawed_model_suite/variant6_compound: 6.0907/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.0907/10 (not FRIES2)
- Human reviewed: no
- Probe flags: empty_file_list, evaluation_contract_v2, high_impact_deployment_claim, high_impact_without_full_disclosure, missing_data_disclosure, missing_ethical_considerations, missing_evaluation, missing_intended_use, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

6.0907 / 10

Dimension	Score
SAFETY	6.9521
FAIRNESS	8.0
INTEGRITY	5.848
ROBUSTNESS	8.6535
EXPLAINABILITY	1.0

Score type: original_FRIES · overall confidence 0.9518

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	/models/flawed_model_suite/variant6_compound
Frozen revision	local

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	8	8	8
ROBUSTNESS	9	9	8
INTEGRITY	5	5	8
EXPLAINABILITY	1	1	1
SAFETY	7	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9473
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	0.7761
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.326535, "ci_upper": 0.412892, "method": "bootstrap_percentile", "point": 0.36831}
eo_ci	{"B": 1000, "ci_lower": 0.110303, "ci_upper": 0.246776, "method": "bootstrap_percentile", "point": 0.1665}
f1_ci	{"B": 1000, "ci_lower": 0.090088, "ci_upper": 0.233899, "method": "bootstrap_percentile", "point": 0.160639}
groups	{"0": {"f1": 0.561485, "fpr": 0.018704, "n": 2461, "positive_rate": 0.065827, "tpr": 0.449814}, "1": {"f1": 0.722124, "fpr": 0.144231, "n": 539, "positive_rate": 0.434137, "tpr": 0.616314}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "/models/flawed_model_suite/variant6_compound", "num_labels": 2, "revision": "local", "task_type": "binary_classification"}
model_ref	/models/flawed_model_suite/variant6_compound
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	local
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.160639
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.1665
demographic_parity_difference	0.36831

Evidence:

Evidence ID	Hash	URI
34850e56-a5ab-46a5-845a-18a999a2181c	sha256:800e8683c396cb9a29663088ab13bbfa4abeb3f4eb8f35ad29e365c756058e0a	s3://trustlens/evidence/53dbc432-fcbc-4a2b-8b4f-4ba7de5c982c/34850e56-a5ab-46a5-845a-18a999a2181c.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	/models/flawed_model_suite/variant6_compound
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.005667, "ci_upper": 0.02, "method": "bootstrap_percentile", "point": 0.013}, "clean_accuracy": {"ci_lower": 0.872741, "ci_upper": 0.895608, "method": "wilson_score", "point": 0.884667}, "robust_accuracy": {"ci_lower": 0.859221, "ci_upper": 0.883162, "method": "wilson_score", "point": 0.871667}}
probe_status	EVALUATED
accuracy_drop	0.013
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.8847
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	local
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.8717
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9853
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.027

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.013, ci_lower=0.005667, ϵ =0.05)
relative_degradation	0.9853
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
4397c1b8-62c0-4e7b-b2e9-df8c5e8a9868	sha256:d5c36be11fcab609590e590c5850e59ee26b8ddb65759f591879a36d3fd1d5bb	s3://trustlens/evidence/53dbc432-fcbc-4a2b-8b4f-4ba7de5c982c/4397c1b8-62c0-4e7b-b2e9-df8c5e8a9868.json

INTEGRITY (confidence 0.7761)

unpinned_revision

empty_file_list

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "1059 chars", "pass": true}, "files_listed": {"detail": "empty or missing files list", "pass": false}, "files_listing_recorded": {"detail": "empty file list \u2014 no listing fingerprint identity", "files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "model_revision": "local", "pass": false}, "license_declared": {"detail": "other", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training_data, evaluation", "pass": true}, "revision_pinned": {"detail": "non-sha-like revision=local", "pass": false}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "files_listing_currency": {"reason": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim.", "status": "INSUFFICIENT EVIDENCE"}, "identity_recording": {"reason": "Hub file manifest missing \u2014 listing-dependent identity claims insufficient", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "files_listing_reverification": "not performed", "hash_comparison": "not performed", "hub_files": [], "live_files_fingerprint": null, "revision": "local", "sha_like": false, "trusted_reference": null, "weight_hash": null}</pre>
disclosure	<pre>{"card_chars": 1059, "card_present": true, "license_structured": "other", "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": ["G-INT-LISTING-UNVERIFIED", "G-INT-HASH-REF-MISSING", "G-INT-HASH-LOCAL-MISSING", "I-INT-HASH-UNVERIFIED"], "gates_passed": false}</pre>
uncertainty	{}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
claim_boundary	<pre>{"cryptographic_identity": "Cryptographic identity was not verified: no trusted reference and/or no local artifact hash. Other Integrity metadata checks are independent of this claim.", "files_listing_currency": "Post-import file-listing currency was not re-verified this run (Hub unreachable, nothing recorded at import time, or re-check not attempted). Other Integrity metadata checks are independent of this claim."}</pre>
scored_risk_id	null
risks_triggered	["I-INT-REV-UNPINNED", "I-INT-MANIFEST-MISSING"]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 0.7761, "data_quality": 0.55, "evidence_completeness": 0.85, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0

Evidence:

Evidence ID	Hash	URI
6cd553d6-d894-4ce5-86ae-c1ab78d824d8	sha256:eb85c423bfaa070551a9d53e1d20d2518912bef85791fb473197c0757bf7e809	s3://trustlens/evidence/53dbc432-fcbc-4a2b-8b4f-4ba7de5c982c/6cd553d6-d894-4ce5-86ae-c1ab78d824d8.json

EXPLAINABILITY (confidence 1.0)

missing_intended_use missing_limitations missing_training_data missing_evaluation missing_ethical_considerations
open_claim_vs_restrictive_license no_limitations_but_production_claim needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	{"documentation_completeness": {"coverage_ratio": 0.0, "detail": "0/5 required sections present", "pass": false, "sections_present": 0, "sections_required": 5}, "documentation_contradiction": {"detail": "open claim vs restrictive license, no limitations but production claim", "pass": false}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "missing or trivial body", "pass": false}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "missing or trivial body", "pass": false}, "training_data": {"detail": "missing or trivial body", "pass": false}}
claims	{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (0/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "keyword-level contradiction flags recorded", "status": "EVALUATED"}}
status	EVALUATED
sections	{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": null, "present": false}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": null, "present": false}, "training_data": {"matched_heading": null, "present": false}}
card_chars	1059
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	{"failed_gates": [], "gates_passed": true}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	{"architecture": {"matched_heading": "Model Description", "present": true}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}
claim_boundary	{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}
contradictions	["open_claim_vs_restrictive_license", "no_limitations_but_production_claim"]
coverage_ratio	0.0
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE", "E-DOC-CONTRADICTION"]
proposed_mapping	False
sections_present	0
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
methodology_version	tl-explainability-v1.0
documentation_source	null

Evidence:

Evidence ID	Hash	URI
7e52df2d-04fe-41d7-a98c-4881e9c9ca5b	sha256:bfaa951256e633906b326f8c90bc3e812e534a6bb582ae499b37b4459c14ac5c	s3://trustlens/evidence/53dbc432-fcbc-4a2b-8b4f-4ba7de5c982c/7e52df2d-04fe-41d7-a98c-4881e9c9ca5b.json

SAFETY (confidence 1.0)

missing_misuse_risks missing_privacy missing_security_warnings missing_data_disclosure high_impact_deployment_claim
high_impact_without_full_disclosure needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	{"data_disclosure": {"detail": "missing or trivial body", "pass": false}, "disclosure_completeness": {"checks_present": 0, "checks_required": 4, "coverage_ratio": 0.0, "detail": "0/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}
claims	{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (0/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "lexical documentation evidence: 1 listed phrase(s) matched", "status": "EVALUATED"}}
status	EVALUATED
behavior	{"benign_fpr": 0.02958333333333333, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.6576402321083172, "harmful_recall": 0.5416666666666667, "min_severe_n": 30, "model_ref": "/models/flawed_model_suite/variant6_compound", "model_revision": "local", "n_evaluated": 3000, "overall_fnr": 0.4583333333333333, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.30141843971631205, "severe_fnr_ci": [0.250823432387957, 0.3573511650114418], "severe_n": 282, "status": "EVALUATED"}
fnr_ratio	0.6576402321083172
card_chars	1059
severe_fnr	0.30141843971631205
limitations	["Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score."]
reliability	{"failed_gates": [], "gates_passed": true}
uncertainty	{}
bonus_checks	{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	0
claim_boundary	{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}
coverage_ratio	0.0
scored_risk_id	null
checks_required	4
required_checks	{"data_disclosure": {"matched_heading": null, "present": false}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}
risks_triggered	["S-GOV-DISCLOSURE-GAP"]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
high_impact_claims	["production_ready"]
methodology_version	tl-safety-v1.0

documentation_source	null
----------------------	------

Evidence:

Evidence ID	Hash	URI
742e0283-0380-4961-9c02-7825eaa6adbc	sha256:1af80b3b9fa356ed218efe7c27e2f84b6d6d4dfc0004a476c22e0e91ba88f6b9	s3://trustlens/evidence/fcbc-4a2b-8b4f-4ba7de5c742e0283-0380-4961-9c02

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.013, ci_lower=0.005667, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

INTEGRITY

Status	EVALUATED
Risk decision	risk detected — risks: I-INT-REV-UNPINNED, I-INT-MANIFEST-MISSING
Gate / rule	G-INT-LISTING-UNVERIFIED, G-INT-HASH-REF-MISSING, G-INT-HASH-LOCAL-MISSING, I-INT-HASH-UNVERIFIED
Human O/S/D	O=5 · S=5 · D=8
FRIES dimension score	5.848
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE, E-DOC-CONTRADICTION
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=1 · S=1 · D=1
FRIES dimension score	1.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=7 · S=6 · D=8
FRIES dimension score	6.952
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	8	8	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.167. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.885, robust_accuracy=0.872, accuracy_drop=0.013. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	5	5	8	0.7761	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=0.50. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	1	1	1	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.00. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	7	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.301, fnr_ratio=0.66. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9518 (not FRIES certainty)

Reproducibility Information

Model reference	/models/flawed_model_suite/variant6_compound
Model revision	local
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

No documentation sources recorded for this model.

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.

- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_disclosure_gap_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:55:21.411518Z · canonical JSON is the source of truth; this PDF is a rendered projection.