

Model	martin-ha/toxic-comment-model
Evaluation ID	1429889a-630e-447a-88e0-6659fb3aa1b4
Status	FINALIZED
Created	2026-10-02T16:03:38.192950Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:20.225816Z

AUTO-FINALIZE (NO HUMAN REVIEW)**Human reviewed:** no**Assessment engine:** llm_v1**Disclosure:** LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment**TrustLens FRIES report for martin-ha/toxic-comment-model: 7.6518/10 (AI-AUTONOMOUS, not human-reviewed)**

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 7.6518/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_ethical_considerations, missing_intended_use, missing_license, missing_misuse_risks, missing_privacy, missing_security_warnings, needs_human_review, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES**7.6518 / 10**

Dimension	Score
SAFETY	6.9521
FAIRNESS	8.6535
INTEGRITY	8.0
ROBUSTNESS	8.6535
EXPLAINABILITY	6.0

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	martin-ha/toxic-comment-model
Frozen revision	9842c08b35a4687e7b211187d676986c8c96256d

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	9	9	8
ROBUSTNESS	9	9	8
INTEGRITY	8	8	8
EXPLAINABILITY	6	6	6
SAFETY	7	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.256141, "ci_upper": 0.339842, "method": "bootstrap_percentile", "point": 0.299028}
eo_ci	{"B": 1000, "ci_lower": 0.033413, "ci_upper": 0.143969, "method": "bootstrap_percentile", "point": 0.06286}
f1_ci	{"B": 1000, "ci_lower": 0.010664, "ci_upper": 0.151844, "method": "bootstrap_percentile", "point": 0.078553}
groups	{"0": {"f1": 0.61215, "fpr": 0.012774, "n": 2461, "positive_rate": 0.064608, "tpr": 0.486989}, "1": {"f1": 0.690702, "fpr": 0.067308, "n": 539, "positive_rate": 0.363636, "tpr": 0.549849}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf", "model_ref": "martin-ha/toxic-comment-model", "num_labels": 2, "revision": "9842c08b35a4687e7b211187d676986c8c96256d", "task_type": "binary_classification"}
model_ref	martin-ha/toxic-comment-model
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	9842c08b35a4687e7b211187d676986c8c96256d
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.078553
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.06286
demographic_parity_difference	0.299028

Evidence:

Evidence ID	Hash	URI
85ac7d94-64b2-49ea-b454-110e3fb35dd3	sha256:c86e5dcbddaccabf0307c0753d0f588a1c41017d4a0cc4fe8c351fbda23222c7	s3://trustlens/evidence/1429889a-630e-447a-88e0-6659fb3aa1b4/85ac7d94-64b2-49ea-b454-110e3fb35dd3.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	martin-ha/toxic-comment-model
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.005, "ci_upper": 0.018667, "method": "bootstrap_percentile", "point": 0.012}, "clean_accuracy": {"ci_lower": 0.878648, "ci_upper": 0.90102, "method": "wilson_score", "point": 0.890333}, "robust_accuracy": {"ci_lower": 0.866149, "ci_upper": 0.88955, "method": "wilson_score", "point": 0.878333}}
probe_status	EVALUATED
accuracy_drop	0.012
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.8903
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	9842c08b35a4687e7b211187d676986c8c96256d
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.8783
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9865
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0257

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.012, ci_lower=0.005, ε=0.05)
relative_degradation	0.9865
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
756c6ecd-d936-4392-9bee-4a707554ff49	sha256:5075a764fe28e214de2728f8d244ad209b74b2ebc1644b6e31c221f60ad2abb1	s3://trustlens/evidence/1429889a-630e-447a-88e0-6659fb3ad936-4392-9bee-4a707554ff49.json

INTEGRITY (confidence 1.0)

missing_license

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "3184 chars", "pass": true}, "files_listed": {"detail": "8 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:429bd407f05bb6ae093d100340f5802ff08ac1becb81b6ffb33586f1c791176c", "model_revision": "9842c08b35a4687e7b211187d676986c8c96256d", "pass": true}, "license_declared": {"detail": "no structured license", "pass": false}, "reproducibility_claims": {"detail": "signal groups: training_data, evaluation", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license check ran; absence recorded as named risk", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:429bd407f05bb6ae093d100340f5802ff08ac1becb81b6ffb33586f1c791176c", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", "README.md", "config.json", "pytorch_model.bin", "special_tokens_map.json", "tokenizer.json", "tokenizer_config.json", "vocab.txt"], "live_files_fingerprint": "sha256:429bd407f05bb6ae093d100340f5802ff08ac1becb81b6ffb33586f1c791176c", "revision": "9842c08b35a4687e7b211187d676986c8c96256d", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf hub lfs", "value": "569aed60978bec9cdc5a90e660fe860e2eccd4f72479c1aac0c9b6c64a581e94"}, "weight_hash": {"algo": "sha256", "value": "569aed60978bec9cdc5a90e660fe860e2eccd4f72479c1aac0c9b6c64a581e94"}}</pre>
disclosure	<pre>{"card_chars": 3184, "card_present": true, "license_structured": null, "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	risk_detected
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>["I-INT-LICENSE-UNDISCLOSED"]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "pytorch_model.bin", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
100e1e11-0887-4cc9-9cab-44cfd0e8e500	sha256:51185a11dc846eaac5e5b14b6b7b4db22d3d6157b4eda1a648902ccbc47c9054	s3://trustlens/evidence/1429889a-630e-447a-88e0

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_ethical_considerations

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	{"documentation_completeness": {"coverage_ratio": 0.6, "detail": "3/5 required sections present", "pass": false, "sections_present": 3, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "Evaluation results", "pass": true}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "Limitations and Bias", "pass": true}, "training_data": {"detail": "Training data", "pass": true}}
claims	{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (3/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}
status	EVALUATED
sections	{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": "Evaluation results", "present": true}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": "Limitations and Bias", "present": true}, "training_data": {"matched_heading": "Training data", "present": true}}
card_chars	3184
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	{"failed_gates": [], "gates_passed": true}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	{"architecture": {"matched_heading": "Model description", "present": true}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": "How to use", "present": true}}
claim_boundary	{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}
contradictions	[]
coverage_ratio	0.6
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	3
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
methodology_version	tl-explainability-v1.0
documentation_source	{"content_length": 3184, "documentation_content_hash": "sha256:65856ba28d1792de2497cdd887a46cdd0353e85e449a5e7a955a54f8b861e5e", "documentation_revision": "9842c08b35a4687e7b211187d676986c8c96256d", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/martin-ha/toxic-comment-model/blob/9842c08b35a4687e7b211187d676986c8c96256d/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "martin-ha/toxic-comment-model", "source_model_revision": "9842c08b35a4687e7b211187d676986c8c96256d"}

Evidence:

Evidence ID	Hash	URI
	sha256:6697ac97e9d52628fab77260fe5f9952823de4b767347b953718b719cf635557	

5f9f5617-
c72c-4a43-
b121-30ced6eeb256

s3://trustlens/evidence/
1429889a-630e-447a-88e0-6659fb3aa1b4/5f9f5
c72c-4a43-b121-30ced6eeb256.json

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "Training data", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.0175, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-1eba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.6227296315516346, "harmful_recall": 0.5216666666666667, "min_severe_n": 30, "model_ref": "martin-ha/toxic-comment-model", "model_revision": "9842c08b35a4687e7b211187d676986c8c96256d", "n_evaluated": 3000, "overall_fnr": 0.4783333333333333, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.2978723404255319, "severe_fnr_ci": [0.24750214520734554, 0.3536755700027568], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.6227296315516346
card_chars	3184
severe_fnr	0.2978723404255319
limitations	[<code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	1
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": "Training data", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	[<code>"S-GOV-DISCLOSURE-GAP"</code>]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	[]
methodology_version	tl-safety-v1.0
documentation_source	<code>{"content_length": 3184, "documentation_content hash": "sha256:65856ba28d1792de2497cdd887a46cdde0353e85e449a5e7a955a54f8b861e5e",</code>

```
"documentation_revision": "9842c08b35a4687e7b211187d676986c8c96256d",
"documentation_source_type": "model_card", "documentation_url": "https://
huggingface.co/martin-ha/toxic-comment-model/blob/
9842c08b35a4687e7b211187d676986c8c96256d/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "martin-ha/toxic-comment-model",
"source_model_revision": "9842c08b35a4687e7b211187d676986c8c96256d"}
```

Evidence:

Evidence ID	Hash	URI
8154d4d1-a24b-4e48-a377-e98133254879	sha256:bf6bb9093208a3d14d2c03bc73df763d3e4fa45e50e2727c61807bd54755a470	s3://trustlens/evidence/1429889a-630e-447a-88e0-6659fb3aa1b4/8154d4d1-a24b-4e48-a377-e98133254879.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.012, ci_lower=0.005, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

INTEGRITY

Status	EVALUATED
Risk decision	risk detected — risks: I-INT-LICENSE-UNDISCLOSED
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.60
Human O/S/D	O=6 · S=6 · D=6
FRIES dimension score	6.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	O=7 · S=6 · D=8
FRIES dimension score	6.952
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	9	9	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.063. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.890, robust_accuracy=0.878, accuracy_drop=0.012. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	8	8	8	1.0	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=0.83. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	6	6	6	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.60. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	7	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.298, fnr_ratio=0.62. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	martin-ha/toxic-comment-model
Model revision	9842c08b35a4687e7b211187d676986c8c96256d
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/martin-ha/toxic-comment-model/blob/9842c08b35a4687e7b211187d676986c8c96256d/README.md	9842c08b35	sha256:65856ba28d1	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:20.225816Z · canonical JSON is the source of truth; this PDF is a rendered projection.