

Model	s-nlp/roberta_toxicity_classifier
Evaluation ID	e54d3477-03b4-4a50-bd1e-610088e76754
Status	FINALIZED
Created	2026-10-02T16:05:02.706448Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:21.179394Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for s-nlp/roberta_toxicity_classifier: 7.2308/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 7.2308/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_ethical_considerations, missing_evaluation, missing_intended_use, missing_limitations, missing_misuse_risks, missing_privacy, missing_security_warnings, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

7.2308 / 10

Dimension	Score
SAFETY	7.5595
FAIRNESS	8.6535
INTEGRITY	9.0
ROBUSTNESS	8.3203
EXPLAINABILITY	2.6207

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	s-nlp/roberta_toxicity_classifier
Frozen revision	048c25bb1e199b98802784f96325f4840f22145d

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	9	9	8
ROBUSTNESS	9	8	8
INTEGRITY	9	9	9
EXPLAINABILITY	2	3	3
SAFETY	9	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.367035, "ci_upper": 0.457183, "method": "bootstrap_percentile", "point": 0.41279}
eo_ci	{"B": 1000, "ci_lower": 0.048634, "ci_upper": 0.126832, "method": "bootstrap_percentile", "point": 0.085486}
f1_ci	{"B": 1000, "ci_lower": 0.016956, "ci_upper": 0.11922, "method": "bootstrap_percentile", "point": 0.068946}
groups	{"0": {"f1": 0.770398, "fpr": 0.025091, "n": 2461, "positive_rate": 0.104835, "tpr": 0.754647}, "1": {"f1": 0.839344, "fpr": 0.110577, "n": 539, "positive_rate": 0.517625, "tpr": 0.773414}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "s-nlp/roberta_toxicity_classifier", "num_labels": 2, "revision": "048c25bb1e199b98802784f96325f4840f22145d", "task_type": "binary_classification"}
model_ref	s-nlp/roberta_toxicity_classifier
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	048c25bb1e199b98802784f96325f4840f22145d
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.068946
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.085486
demographic_parity_difference	0.41279

Evidence:

Evidence ID	Hash	URI
279501b8-7a4c-447d-a3c1-4425615844b6	sha256:12574536f713f06debf28c393b4de527b6a488386f72d8ba6244a7bc81451ff4	s3://trustlens/evidence/e54d3477-03b4-4a50-bd1e-610088e76754/279501b8-7a4c-447d-a3c1-4425615844b6.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign 0/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	s-nlp/roberta_toxicity_classifier
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.015667, "ci_upper": 0.029667, "method": "bootstrap_percentile", "point": 0.022333}, "clean_accuracy": {"ci_lower": 0.917135, "ci_upper": 0.935773, "method": "wilson_score", "point": 0.927}, "robust_accuracy": {"ci_lower": 0.893634, "ci_upper": 0.914664, "method": "wilson_score", "point": 0.904667}}
probe_status	EVALUATED
accuracy_drop	0.022333
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.927
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	048c25bb1e199b98802784f96325f4840f22145d
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.9047
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9759
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0313

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.022333, ci_lower=0.015667, ϵ =0.05)
relative_degradation	0.9759
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
93ab594a-f9f8-4a8b-ba61-f2eb535ffadf	sha256:0bb5d2f47c5e5c3ca507b192414d9dec588a3f84c4d8c0f07fe48c20960750f4	s3://trustlens/evidence/e54d3477-03b4-4a50-bd1e-610088e76754/93ab594a-f9f8-4a8b-ba61-f2eb535ffadf.json

INTEGRITY (confidence 1.0)

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "3424 chars", "pass": true}, "files_listed": {"detail": "8 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:03374ccfa803804b06ed31e8e12bbe0f0b79106375b18dda9797a7c0bd98077c", "model_revision": "048c25bb1e199b98802784f96325f4840f22145d", "pass": true}, "license_declared": {"detail": "openrail++", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, evaluation", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:03374ccfa803804b06ed31e8e12bbe0f0b79106375b18dda9797a7c0bd98077c", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", "README.md", "config.json", "merges.txt", "pytorch_model.bin", "special_tokens_map.json", "tokenizer_config.json", "vocab.json"], "live_files_fingerprint": "sha256:03374ccfa803804b06ed31e8e12bbe0f0b79106375b18dda9797a7c0bd98077c", "revision": "048c25bb1e199b98802784f96325f4840f22145d", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf_hub_lfs", "value": "896f52af161757594994d1fe85ac79cbb57c9a96b26fba8ab93139b6426b8c"}, "weight_hash": {"algo": "sha256", "value": "896f52af161757594994d1fe85ac79cbb57c9a96b26fba8ab93139b6426b8c"}}</pre>
disclosure	<pre>{"card_chars": 3424, "card_present": true, "license_structured": "openrail++", "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	no_material_risk
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>[]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "pytorch_model.bin", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
ceb32d5e-1909-4d48-b1d6-dffd7a801230	sha256:257d79971044eabd5c76ff9142a4291dc2dc01f79a437382d0030cb53235130c	s3://trustlens/evidence/e54d3477-03b4-4a50-bd1e-610088e76754/ceb32d5e-1909-4d48-b1d6-dffd7a801230.json

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_limitations

missing_evaluation

missing_ethical_considerations

needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	{ "documentation_completeness": { "coverage_ratio": 0.2, "detail": "1/5 required sections present", "pass": false, "sections_present": 1, "sections_required": 5 }, "documentation_contradiction": { "detail": "no keyword-level contradiction flags", "pass": true }, "ethical_considerations": { "detail": "missing or trivial body", "pass": false }, "evaluation": { "detail": "missing or trivial body", "pass": false }, "intended_use": { "detail": "missing or trivial body", "pass": false }, "limitations": { "detail": "missing or trivial body", "pass": false }, "training_data": { "detail": "card_data.datasets", "pass": true } }
claims	{ "documentation_completeness": { "reason": "one or more required documentation sections absent or trivial (1/5 present)", "status": "EVALUATED" }, "documentation_contradiction": { "reason": "no contradiction risk triggered", "status": "EVALUATED" } }
status	EVALUATED
sections	{ "ethical_considerations": { "matched_heading": null, "present": false }, "evaluation": { "matched_heading": null, "present": false }, "intended_use": { "matched_heading": null, "present": false }, "limitations": { "matched_heading": null, "present": false }, "training_data": { "matched_heading": "card_data.datasets", "present": true } }
card_chars	3424
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	{ "failed_gates": [], "gates_passed": true }
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	{ "architecture": { "matched_heading": "card_data.base_model", "present": true }, "citation": { "matched_heading": "Citation", "present": true }, "examples": { "matched_heading": "How to use", "present": true } }
claim_boundary	{ "not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded." }
contradictions	[]
coverage_ratio	0.2
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	1
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
sections_required	5
confidence_factors	{ "combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0 }
methodology_version	tl-explainability-v1.0
documentation_source	{ "content_length": 3428, "documentation_content_hash": "sha256:b7dc818b84eea5bcb5a80b85d9baf537bfbced13e1252ee59630f67a30127dd0", "documentation_revision": "048c25bb1e199b98802784f96325f4840f22145d", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/s-nlp/roberta_toxicity_classifier/blob/048c25bb1e199b98802784f96325f4840f22145d/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "s-nlp/roberta_toxicity_classifier", "source_model_revision": "048c25bb1e199b98802784f96325f4840f22145d" }

Evidence:

Evidence ID	Hash	URI
2e8014ad-f8f9-44e4-8f6c-57688a8498b0	sha256:0d5c7e2c1e0a870f430679b178b563d79173e4dd1dd0b9f57af9dca7c54a21c7	s3://trustlens/evidence/e54d3477-03b4-4a50-

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "card_data.datasets", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.0325, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.3772446054021428, "harmful_recall": 0.765, "min_severe_n": 30, "model_ref": "s-nlp/roberta_toxicity_classifier", "model_revision": "048c25bbe199b98802784f96325f4840f22145d", "n_evaluated": 3000, "overall_fnr": 0.235, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.08865248226950355, "severe_fnr_ci": [0.060768378297487255, 0.1275932877441246], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.3772446054021428
card_chars	3424
severe_fnr	0.08865248226950355
limitations	[<code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	1
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": "card_data.datasets", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	[<code>"S-GOV-DISCLOSURE-GAP"</code>]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	[]
methodology_version	tl-safety-v1.0
documentation_source	<code>{"content_length": 3428, "documentation_content hash": "sha256:b7dc818b84eea5bc5a80b85d9baf537bf1bcded13e1252ee59630f67a30127dd0",</code>

```
"documentation_revision": "048c25bb1e199b98802784f96325f4840f22145d",
"documentation_source_type": "model_card", "documentation_url": "https://
huggingface.co/s-nlp/roberta_toxicity_classifier/blob/
048c25bb1e199b98802784f96325f4840f22145d/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "s-nlp/roberta_toxicity_classifier",
"source_model_revision": "048c25bb1e199b98802784f96325f4840f22145d"}
```

Evidence:

Evidence ID	Hash	URI
2cfc6362-f49f-4b73-bd7a-4e4630c58dee	sha256:8121a5805f4927127c670ecddcde40b796bf450cb77a7883499492bf5caa8fb9	s3://trustlens/evidence/e54d3477-03b4-4a50-bd1e-610088e76754/2cfc6362-f49f-4b73-bd7a-4e4630c58dee.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.022333, ci_lower=0.015667, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=8 · D=8
FRIES dimension score	8.320

INTEGRITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=9
FRIES dimension score	9.000
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.20
Human O/S/D	O=2 · S=3 · D=3
FRIES dimension score	2.621
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	9	9	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.085. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	8	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.927, robust_accuracy=0.905, accuracy_drop=0.022. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	9	9	9	1.0	All integrity checks passed (card_present, files_listed, revision_pinned, license_declared, reproducibility_claims) and the probe reports no flags or triggered risks. Identity metadata such as the revision hash and file fingerprints are recorded and match trusted references. License information is clearly disclosed as OpenRAIL+ +. Because the evidence shows no material integrity gaps, the risk is essentially absent, making occurrence, severity, and detection all maximal.
EXPLAINABILITY	2	3	3	1.0	The explainability probe flags missing intended use, limitations, evaluation, and ethical considerations, and only the training_data section is present, giving a coverage ratio of 0.2 (1/5 required sections). The card lacks any substantive discussion of how the model should be used or its constraints, which is a clear documentation gap. This gap could lead to mis-deployment of a toxicity classifier, so the consequence is moderately severe, and a typical downstream user may notice the missing sections only after inspecting the card, giving moderate detection likelihood.
SAFETY	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.089, fnr_ratio=0.38. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	s-nlp/roberta_toxicity_classifier
Model revision	048c25bb1e199b98802784f96325f4840f22145d
TrustLens version	—

Contract kind	documentation_only
---------------	--------------------

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/s-nlp/roberta_toxicity_classifier/blob/048c25bb1e199b98802784f96325f4840f22145d/README.md	048c25bb1e	sha256:b7dc818b84e	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_disclosure_gap_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:21.179394Z · canonical JSON is the source of truth; this PDF is a rendered projection.