

|               |                                                               |
|---------------|---------------------------------------------------------------|
| Model         | BunnyNoBugs/rubert-tiny2-russe-toxicity                       |
| Evaluation ID | 9984a269 - fd93 - 4a29 - a3d4 - 63ca70bbb9ff                  |
| Status        | FINALIZED                                                     |
| Created       | 2026-10-02T16:08:46.949320Z                                   |
| Report        | report_v1 · version 1 · generated 2026-10-02T17:54:22.846958Z |

**AUTO-FINALIZE (NO HUMAN REVIEW)****Human reviewed:** no**Assessment engine:** llm\_v1**Disclosure:** LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY\_HEURISTIC\_OSD\_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

**Executive Assessment****TrustLens FRIES report for BunnyNoBugs/rubert-tiny2-russe-toxicity: 6.1999/10 (AI-AUTONOMOUS, not human-reviewed)**

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.1999/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation\_contract\_v2, missing\_data\_disclosure, missing\_ethical\_considerations, missing\_intended\_use, missing\_limitations, missing\_misuse\_risks, missing\_privacy, missing\_security\_warnings, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

**FRIES****6.1999 / 10**

| Dimension      | Score  |
|----------------|--------|
| SAFETY         | 4.5789 |
| FAIRNESS       | 8.6535 |
| INTEGRITY      | 8.6535 |
| ROBUSTNESS     | 7.1138 |
| EXPLAINABILITY | 2.0    |

Score type: original\_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

**Evaluation Contract**

|                               |                    |
|-------------------------------|--------------------|
| Contract kind                 | documentation_only |
| Task                          | —                  |
| Dataset (comparability label) | —                  |
| Config                        | —                  |
| Pairing ID                    | —                  |

**Model Identity**

|                 |                                          |
|-----------------|------------------------------------------|
| Model reference | BunnyNoBugs/rubert-tiny2-russe-toxicity  |
| Frozen revision | b136cda1a1417bec27c492d0fe1eaa3f878606e4 |

Dataset Identity

|                      |   |
|----------------------|---|
| Dataset key          | — |
| Dataset revision     | — |
| User dataset ID      | — |
| Dataset content hash | — |

Local Execution Environment

|                   |                                          |
|-------------------|------------------------------------------|
| Execution device  | GPU (NVIDIA GeForce RTX 4060 Laptop GPU) |
| Inference backend | local_hf                                 |
| CUDA available    | True                                     |

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

| Aspect         | O (Occurrence) | S (Severity) | D (Detection) |
|----------------|----------------|--------------|---------------|
| FAIRNESS       | 9              | 9            | 8             |
| ROBUSTNESS     | 9              | 5            | 8             |
| INTEGRITY      | 9              | 9            | 8             |
| EXPLAINABILITY | 2              | 2            | 2             |
| SAFETY         | 2              | 6            | 8             |

Source: osd\_agent\_autonomous · human reviewed: no

Confidence summary

Method: geometric\_mean\_v1 · proposed calibration: True

| Scope          | Confidence |
|----------------|------------|
| Overall        | 0.9966     |
| FAIRNESS       | 0.983      |
| ROBUSTNESS     | 1.0        |
| INTEGRITY      | 1.0        |
| EXPLAINABILITY | 1.0        |
| SAFETY         | 1.0        |

Evidence strength only — not correctness or O/S/D



**FAIRNESS (confidence 0.983)**

evaluation\_contract\_v2

user\_defined\_local\_dataset

wide\_ci\_f1

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| seed                 | 42                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| dp_ci                | {"B": 1000, "ci_lower": 0.08104, "ci_upper": 0.154764, "method": "bootstrap_percentile", "point": 0.119371}                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| eo_ci                | {"B": 1000, "ci_lower": 0.054115, "ci_upper": 0.158253, "method": "bootstrap_percentile", "point": 0.097628}                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| f1_ci                | {"B": 1000, "ci_lower": 0.088295, "ci_upper": 0.241554, "method": "bootstrap_percentile", "point": 0.165123}                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| groups               | {"0": {"f1": 0.18677, "fpr": 0.089872, "n": 2461, "positive_rate": 0.099553, "tpr": 0.178439}, "1": {"f1": 0.351893, "fpr": 0.1875, "n": 539, "positive_rate": 0.218924, "tpr": 0.238671}}                                                                                                                                                                                                                                                                                                                                                                |
| inference            | {"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "BunnyNoBugs/rubert-tiny2-russe-toxicity", "num_labels": 2, "revision": "b136cd1a1417bec27c492d0fe1eaa3f878606e4", "task_type": "binary_classification"} |
| model_ref            | BunnyNoBugs/rubert-tiny2-russe-toxicity                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| predictor            | inference_backend                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| excess_dpd           | 0.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| dataset_key          | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| min_group_n          | 30                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| n_evaluated          | 3000                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| text_column          | text                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| group_column         | identity_ref                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| missing_rows         | {"missing_count": 0, "missing_reasons": {}}                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| excluded_rows        | {"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| fairness_mode        | user_defined_local                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| label_mapping        | [{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| target_column        | label                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| dataset_format       | csv                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| label_encoding       | {"0": 0, "1": 1}                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| label_rate_gap       | 0.504795                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| model_revision       | b136cd1a1417bec27c492d0fe1eaa3f878606e4                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| excluded_groups      | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| observed_groups      | [{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| dataset_revision     | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| evaluation_class     | user_dataset_v2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| dataset_row_count    | 3000                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| confidence_factors   | {"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| dataset_content_id   | c5d8778e-1eba-4a8c-a55c-82487c90a1b2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| inference_executed   | True                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| needs_human_review   | True                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| subgroup_f1_spread   | 0.165123                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| probe_status_reason  | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| sensitive_attribute  | identity_ref                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| min_group_n_observed | 539                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| positive_label_index | 1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |

|                               |          |
|-------------------------------|----------|
| equalized_odds_difference     | 0.097628 |
| demographic_parity_difference | 0.119371 |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| 64cd79f8-d2c4-468f-8aa3-548ae0650ea3 | sha256:85c6672a18dde96bb97b77c22f72f8ba2440e4b811c5d834da61fd11e1891e20 | s3://trustlens/evidence/9984a269-fd93-4a29-a3d4-63ca70bbb9ff/64cd79f8-d2c4-468f-8aa3-548ae0650ea3.json |

**ROBUSTNESS (confidence 1.0)**

evaluation\_contract\_v2

user\_defined\_local\_dataset

| Metric              | Value                                                                                                                                                                                                                                                                                                                                                |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| norm                | discrete                                                                                                                                                                                                                                                                                                                                             |
| note                | Layer A evidence only – R-ROB-PERT detection does not assign O/S/D                                                                                                                                                                                                                                                                                   |
| seed                | 42                                                                                                                                                                                                                                                                                                                                                   |
| attack              | char_swap                                                                                                                                                                                                                                                                                                                                            |
| status              | EVALUATED                                                                                                                                                                                                                                                                                                                                            |
| dataset             | {"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}                                                                                                                                                                                                                                                             |
| epsilon             | 0.03                                                                                                                                                                                                                                                                                                                                                 |
| inference           | {"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}             |
| model_ref           | BunnyNoBugs/rubert-tiny2-russe-toxicity                                                                                                                                                                                                                                                                                                              |
| n_samples           | 3000                                                                                                                                                                                                                                                                                                                                                 |
| task_type           | null                                                                                                                                                                                                                                                                                                                                                 |
| dataset_key         | null                                                                                                                                                                                                                                                                                                                                                 |
| limitations         | []                                                                                                                                                                                                                                                                                                                                                   |
| max_changes         | 3                                                                                                                                                                                                                                                                                                                                                    |
| n_evaluated         | 3000                                                                                                                                                                                                                                                                                                                                                 |
| n_requested         | 3000                                                                                                                                                                                                                                                                                                                                                 |
| reliability         | {"failed_gates": [], "gates_passed": true}                                                                                                                                                                                                                                                                                                           |
| text_column         | text                                                                                                                                                                                                                                                                                                                                                 |
| uncertainty         | {"accuracy_drop": {"B": 1000, "ci_lower": 0.035, "ci_upper": 0.057667, "method": "bootstrap_percentile", "point": 0.046}, "clean_accuracy": {"ci_lower": 0.748133, "ci_upper": 0.778526, "method": "wilson_score", "point": 0.763667}, "robust_accuracy": {"ci_lower": 0.701288, "ci_upper": 0.733488, "method": "wilson_score", "point": 0.717667}} |
| probe_status        | EVALUATED                                                                                                                                                                                                                                                                                                                                            |
| accuracy_drop       | 0.046                                                                                                                                                                                                                                                                                                                                                |
| excluded_rows       | {"missing_target": 0, "missing_text": 0, "unmapped_label": 0}                                                                                                                                                                                                                                                                                        |
| label_mapping       | [{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]                                                                                                                                                                                                                                                     |
| target_column       | label                                                                                                                                                                                                                                                                                                                                                |
| aspect_scoring      | no_material_risk                                                                                                                                                                                                                                                                                                                                     |
| clean_accuracy      | 0.7637                                                                                                                                                                                                                                                                                                                                               |
| dataset_format      | csv                                                                                                                                                                                                                                                                                                                                                  |
| label_encoding      | {"0": 0, "1": 1}                                                                                                                                                                                                                                                                                                                                     |
| model_revision      | b136cd1a1417bec27c492d0fe1eaa3f878606e4                                                                                                                                                                                                                                                                                                              |
| scored_risk_id      | null                                                                                                                                                                                                                                                                                                                                                 |
| risks_triggered     | []                                                                                                                                                                                                                                                                                                                                                   |
| robust_accuracy     | 0.7177                                                                                                                                                                                                                                                                                                                                               |
| dataset_revision    | null                                                                                                                                                                                                                                                                                                                                                 |
| evaluation_class    | user_dataset_v2                                                                                                                                                                                                                                                                                                                                      |
| n_perturb_failed    | 0                                                                                                                                                                                                                                                                                                                                                    |
| proposed_mapping    | False                                                                                                                                                                                                                                                                                                                                                |
| degradation_ratio   | 0.9398                                                                                                                                                                                                                                                                                                                                               |
| confidence_factors  | {"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}                                                                                                                                                                                                                                                       |
| dataset_content_id  | c5d8778e-1eba-4a8c-a55c-82487c90a1b2                                                                                                                                                                                                                                                                                                                 |
| inference_executed  | True                                                                                                                                                                                                                                                                                                                                                 |
| n_label_compatible  | 3000                                                                                                                                                                                                                                                                                                                                                 |
| attack_success_rate | 0.074                                                                                                                                                                                                                                                                                                                                                |

|                          |                                                                                        |
|--------------------------|----------------------------------------------------------------------------------------|
| methodology_version      | tl-methodology-v1.0                                                                    |
| probe_status_reason      | no material perturbation risk at 95% CI (drop=0.046, ci_lower=0.035, $\epsilon$ =0.05) |
| relative_degradation     | 0.9398                                                                                 |
| label_compat_fraction    | 1.0                                                                                    |
| perturbation_coverage    | 1.0                                                                                    |
| n_successfully_perturbed | 3000                                                                                   |

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|------------------------------------------------------------------------|
| d9b343dd-374d-46ec-8b84-96c31fe8b13a | sha256:fb28e7fe42ca04bb4963935cce26dc1db46a598d941d371bdd4a8890bec3275f | s3://trustlens/evidence/fd93-4a29-a3d4-63ca70bbd9b343dd-374d-46ec-8b84 |

INTEGRITY (confidence 1.0)

| Metric                | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                  | Layer A evidence only – Integrity risks do not assign O/S/D                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| checks                | <pre>{"card_present": {"detail": "6593 chars", "pass": true}, "files_listed": {"detail": "17 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:6a719678bd6c885c56a4d67728baa28d1000989151eee2dbe2abd791c16d580a", "model_revision": "b136cdala1417bec27c492d0feleaa3f878606e4", "pass": true}, "license_declared": {"detail": "mit", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, seed, evaluation, hyperparameters", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| claims                | <pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| status                | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| identity              | <pre>{"files_listing_fingerprint": "sha256:6a719678bd6c885c56a4d67728baa28d1000989151eee2dbe2abd791c16d580a", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", ".gitignore", "README.md", "config.json", "pytorch_model.bin", "runs/Apr24_16-49-49_917c996ec89e/1682355002.9907618/events.out.tfevents.1682355002.917c996ec89e.975.1", "runs/Apr24_16-49-49_917c996ec89e/events.out.tfevents.1682355002.917c996ec89e.975.0", "runs/Apr24_16-59-22_a6fbd89319ff/1682355706.8443892/events.out.tfevents.1682355706.a6fbd89319ff.153.1", "runs/Apr24_16-59-22_a6fbd89319ff/events.out.tfevents.1682355706.a6fbd89319ff.153.0", "runs/Apr24_16-59-22_a6fbd89319ff/events.out.tfevents.1682356038.a6fbd89319ff.153.2", "runs/Apr24_17-13-06_a6fbd89319ff/1682356391.867624/events.out.tfevents.1682356391.a6fbd89319ff.5571.1", "runs/Apr24_17-13-06_a6fbd89319ff/events.out.tfevents.1682356391.a6fbd89319ff.5571.0", "special_tokens_map.json", "tokenizer.json", "tokenizer_config.json", "training_args.bin", "vocab.txt"], "live_files_fingerprint": "sha256:6a719678bd6c885c56a4d67728baa28d1000989151eee2dbe2abd791c16d580a", "revision": "b136cdala1417bec27c492d0feleaa3f878606e4", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf hub lfs", "value": "bffc45b6de2d25c86f643a30ad3ea49c96fbc73bf4f1ed502ecf1356e75a6838"}, "weight_hash": {"algo": "sha256", "value": "bffc45b6de2d25c86f643a30ad3ea49c96fbc73bf4f1ed502ecf1356e75a6838"}}</pre> |
| disclosure            | <pre>{"card_chars": 6593, "card_present": true, "license_structured": "mit", "repro_groups_found": ["training_data", "seed", "evaluation", "hyperparameters"]}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| limitations           | <p>["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]</p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| reliability           | <pre>{"failed_gates": [], "gates_passed": true}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| uncertainty           | <pre>{}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| probe_status          | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| osd_proposals         | <pre>[]</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| aspect_scoring        | no_material_risk                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| claim_boundary        | <pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| scored_risk_id        | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| risks_triggered       | <pre>[]</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| proposed_mapping      | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| methodology_basis     | TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| confidence_factors    | <pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| methodology_version   | tl-integrity-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| artifact_verification | <pre>{"file": "pytorch_model.bin", "performed": true}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |

Evidence:

| Evidence ID | Hash | URI |
|-------------|------|-----|
|-------------|------|-----|

|                                              |                                                                         |                                                                                                                        |
|----------------------------------------------|-------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|
| 10590cd6-<br>eb14-4e70-<br>ae5c-241e5ca5d81a | sha256:7a68c467fff2c8a62b357472112dd96a85d2ee65b4d8cc8f2074f9394c93506f | s3://trustlens/evidence/<br>9984a269-fd93-4a29-<br>a3d4-63ca70bbb9ff/10590cd6-<br>eb14-4e70-<br>ae5c-241e5ca5d81a.json |
|----------------------------------------------|-------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|

**EXPLAINABILITY (confidence 1.0)**

missing\_intended\_use

missing\_limitations

missing\_training\_data

missing\_ethical\_considerations

needs\_human\_review

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Layer A documentation/transparency evidence only – not O/S/D or interpretability quality                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| checks               | <pre>{"documentation_completeness": {"coverage_ratio": 0.2, "detail": "1/5 required sections present", "pass": false, "sections_present": 1, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "Training results", "pass": true}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "missing or trivial body", "pass": false}, "training_data": {"detail": "missing or trivial body", "pass": false}}</pre> |
| claims               | <pre>{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (1/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</pre>                                                                                                                                                                                                                                                                                                                                                                           |
| status               | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| sections             | <pre>{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": "Training results", "present": true}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": null, "present": false}, "training_data": {"matched_heading": null, "present": false}}</pre>                                                                                                                                                                                                                                                                                          |
| card_chars           | 6593                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| limitations          | <p>["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]</p>                                                  |
| reliability          | <pre>{"failed_gates": [], "gates_passed": true}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| osd_proposals        | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| aspect_scoring       | risk_detected                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| bonus_sections       | <pre>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| claim_boundary       | <pre>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</pre>                                                                                                                                                                                         |
| contradictions       | []                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| coverage_ratio       | 0.2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| scored_risk_id       | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| risks_triggered      | ["E-DOC-INCOMPLETE"]                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| sections_present     | 1                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| methodology_basis    | TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| sections_required    | 5                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| confidence_factors   | <pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| methodology_version  | tl-explainability-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| documentation_source | <pre>{"content_length": 6594, "documentation_content_hash": "sha256:f144b7b2461ee82aeb49c16da5bd0585aab890eaf5fb44c11ebcbeb953c86ddc", "documentation_revision": "b136cdala1417bec27c492d0fe1eaa3f878606e4", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/BunnyNoBugs/rubert-tiny2-russe-toxicity/blob/b136cdala1417bec27c492d0fe1eaa3f878606e4/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "BunnyNoBugs/rubert-tiny2-russe-toxicity", "source_model_revision": "b136cdala1417bec27c492d0fe1eaa3f878606e4"}</pre>                                        |

Evidence:

| Evidence ID | Hash                                                                   | URI                                         |
|-------------|------------------------------------------------------------------------|---------------------------------------------|
|             | sha256:cfc049eabc77a13de68374610da00488401c4a8bf19e4a753d663bd714836b6 | s3://trustlens/evidence/9984a269-fd93-4a29- |

19450ec5-  
f16c-4cde-838a-  
e924c35a0c4f

a3d4-63ca70bbb9ff/19450ec5-  
f16c-4cde-838a-  
e924c35a0c4f.json

**SAFETY (confidence 1.0)**

missing\_misuse\_risks

missing\_privacy

missing\_security\_warnings

missing\_data\_disclosure

needs\_human\_review

| Metric               | Value                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| note                 | Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| checks               | <code>{"data_disclosure": {"detail": "missing or trivial body", "pass": false}, "disclosure_completeness": {"checks_present": 0, "checks_required": 4, "coverage_ratio": 0.0, "detail": "0/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>                                                                                                                                                                                                                                                                                                                                      |
| claims               | <code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (0/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| status               | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| behavior             | <code>{"benign_fpr": 0.09833333333333333, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.9716162115964194, "harmful_recall": 0.21166666666666667, "min_severe_n": 30, "model_ref": "BunnyNoBugs/rubert-tiny2-russe-toxicity", "model_revision": "b136cdala1a17bec27c492d0fe1eaa3f878606e4", "n_evaluated": 3000, "overall_fnr": 0.7883333333333333, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.7659574468085106, "severe_fnr_ci": [0.7131687549349371, 0.8115974092622965], "severe_n": 282, "status": "EVALUATED"}</code>                                                                                                           |
| fnr_ratio            | 0.9716162115964194                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| card_chars           | 6593                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| severe_fnr           | 0.7659574468085106                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| limitations          | [ <code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code> ] |
| reliability          | <code>{"failed_gates": [], "gates_passed": true}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| uncertainty          | <code>{}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| bonus_checks         | <code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| probe_status         | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| osd_proposals        | <code>[]</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| aspect_scoring       | risk_detected                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| checks_present       | 0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| claim_boundary       | <code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>                                                                                                                                                 |
| coverage_ratio       | 0.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| scored_risk_id       | null                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| checks_required      | 4                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| required_checks      | <code>{"data_disclosure": {"matched_heading": null, "present": false}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| risks_triggered      | <code>["S-GOV-DISCLOSURE-GAP"]</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| proposed_mapping     | False                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| methodology_basis    | TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| confidence_factors   | <code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| high_impact_claims   | <code>[]</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| methodology_version  | tl-safety-v1.0                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| documentation_source | <code>{"content_length": 6594, "documentation_content_hash": "sha256:f144b7b2461ee82aeb49c16da5bd0585aab890eaf5fb44c11ebcbeb953c86ddc",</code>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |

```
"documentation_revision": "b136cdala1417bec27c492d0fe1eaa3f878606e4",
"documentation_source_type": "model_card", "documentation_url": "https://
huggingface.co/BunnyNoBugs/rubert-tiny2-russe-toxicity/blob/
b136cdala1417bec27c492d0fe1eaa3f878606e4/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "BunnyNoBugs/rubert-tiny2-russe-
toxicity", "source_model_revision": "b136cdala1417bec27c492d0fe1eaa3f878606e4"}
```

Evidence:

| Evidence ID                          | Hash                                                                    | URI                                                                                                    |
|--------------------------------------|-------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| 5e7cba0a-69a3-4a91-80c1-dadfdb9db441 | sha256:915be24762a86176b3d8112a4c0f681b5f335096bf4279abb7aaece69744188c | s3://trustlens/evidence/9984a269-fd93-4a29-a3d4-63ca70bbb9ff/5e7cba0a-69a3-4a91-80c1-dadfdb9db441.json |

## Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

### FAIRNESS

|                              |                           |
|------------------------------|---------------------------|
| <b>Status</b>                | EVALUATED                 |
| <b>Risk decision</b>         | Not scored                |
| <b>Gate / rule</b>           | No gate failures recorded |
| <b>Human O/S/D</b>           | O=9 · S=9 · D=8           |
| <b>FRIES dimension score</b> | 8.653                     |

### ROBUSTNESS

|                              |                                                                                          |
|------------------------------|------------------------------------------------------------------------------------------|
| <b>Status</b>                | EVALUATED — no material perturbation risk at 95% CI (drop=0.046, ci_lower=0.035, ε=0.05) |
| <b>Risk decision</b>         | no material risk                                                                         |
| <b>Gate / rule</b>           | No gate failures recorded                                                                |
| <b>Human O/S/D</b>           | O=9 · S=5 · D=8                                                                          |
| <b>FRIES dimension score</b> | 7.114                                                                                    |

### INTEGRITY

|                              |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Status</b>                | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| <b>Risk decision</b>         | no material risk                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| <b>Gate / rule</b>           | No gate failures recorded                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| <b>Human O/S/D</b>           | O=9 · S=9 · D=8                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| <b>FRIES dimension score</b> | 8.653                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| <b>Limitations</b>           | Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files. |

### EXPLAINABILITY

|                               |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|-------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Status</b>                 | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Risk decision</b>          | risk detected — risks: E-DOC-INCOMPLETE                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| <b>Gate / rule</b>            | No gate failures recorded                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Documentation coverage</b> | 0.20                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| <b>Human O/S/D</b>            | O=2 · S=2 · D=2                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
| <b>FRIES dimension score</b>  | 2.000                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| <b>Limitations</b>            | Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score. |

**SAFETY**

|                        |                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Status                 | EVALUATED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Risk decision          | risk detected — risks: S-GOV-DISCLOSURE-GAP                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
| Gate / rule            | No gate failures recorded                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| Documentation coverage | 0.00                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| Human O/S/D            | O=2 · S=6 · D=8                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| FRIES dimension score  | 4.579                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
| Limitations            | Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score. |

**O/S/D representation**

Methodology status: **LEGACY\_HEURISTIC\_OSD\_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

| Aspect         | O (Occurrence) | S (Severity) | D (Detection) | Evidence strength | Rationale                                                                                                                                                                 |
|----------------|----------------|--------------|---------------|-------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| FAIRNESS       | 9              | 9            | 8             | 0.983             | [PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.098. Heuristic metric-to-O/S/D mapping — not validated science.                                       |
| ROBUSTNESS     | 9              | 5            | 8             | 1.0               | [PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.764, robust_accuracy=0.718, accuracy_drop=0.046. Heuristic metric-to-O/S/D mapping — not validated science. |
| INTEGRITY      | 9              | 9            | 8             | 1.0               | [PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=1.00. Heuristic metric-to-O/S/D mapping — not validated science.                                     |
| EXPLAINABILITY | 2              | 2            | 2             | 1.0               | [PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.20. Heuristic metric-to-O/S/D mapping — not validated science.                                     |
| SAFETY         | 2              | 6            | 8             | 1.0               | [PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.766, fnr_ratio=0.97. Heuristic metric-to-O/S/D mapping — not validated science.                                     |

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

**Reproducibility Information**

|                   |                                          |
|-------------------|------------------------------------------|
| Model reference   | BunnyNoBugs/rubert-tiny2-russe-toxicity  |
| Model revision    | b136cda1a1417bec27c492d0fe1eaa3f878606e4 |
| TrustLens version | —                                        |
| Contract kind     | documentation_only                       |

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

**Documentation Sources**

| Source           | Type       | URL                                                                                                                                                                                                                                                         | Revision   | Hash               | Retrieval |
|------------------|------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------|--------------------|-----------|
| HF card (pinned) | model_card | <a href="https://huggingface.co/BunnyNoBugs/rubert-tiny2-russe-toxicity/blob/b136cda1a1417bec27c492d0fe1eaa3f878606e4/README.md">https://huggingface.co/BunnyNoBugs/rubert-tiny2-russe-toxicity/blob/b136cda1a1417bec27c492d0fe1eaa3f878606e4/README.md</a> | b136cda1a1 | sha256:f144b7b2461 | ok        |

## Limitations

---

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-\* values are evidence-layer detections only; scored\_risk\_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored\_risk\_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect\_scoring=risk\_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect\_scoring=governance\_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage\_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage\_ratio counts present required sections; it is not an explainability quality score.
- high\_impact\_claims are documentation metadata flags; they are never promoted to risks\_triggered and do not map to O/S/D or FRIES.

report\_v1 · report version 1 · generated 2026-10-02T17:54:22.846958Z · canonical JSON is the source of truth; this PDF is a rendered projection.