

Model	JungleLee/bert-toxic-comment-classification
Evaluation ID	5f5c23d1-496e-4ccb-ae94-67ae9624ccd6
Status	FINALIZED
Created	2026-10-02T16:06:32.666868Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:23.834742Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for JungleLee/bert-toxic-comment-classification: 7.1545/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 7.1545/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_ethical_considerations, missing_intended_use, missing_limitations, missing_misuse_risks, missing_privacy, missing_security_warnings, needs_human_review, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

7.1545 / 10

Dimension	Score
SAFETY	7.5595
FAIRNESS	8.0
INTEGRITY	8.6535
ROBUSTNESS	7.5595
EXPLAINABILITY	4.0

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	JungleLee/bert-toxic-comment-classification
Frozen revision	9c59104291a49a877078424cbfd9f3af23c26ef9

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	8	8	8
ROBUSTNESS	9	6	8
INTEGRITY	9	9	8
EXPLAINABILITY	4	4	4
SAFETY	9	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.413323, "ci_upper": 0.495758, "method": "bootstrap_percentile", "point": 0.453158}
eo_ci	{"B": 1000, "ci_lower": 0.150598, "ci_upper": 0.276939, "method": "bootstrap_percentile", "point": 0.216557}
f1_ci	{"B": 1000, "ci_lower": 0.222368, "ci_upper": 0.327519, "method": "bootstrap_percentile", "point": 0.273143}
groups	{"0": {"f1": 0.540466, "fpr": 0.119982, "n": 2461, "positive_rate": 0.186916, "tpr": 0.732342}, "1": {"f1": 0.813609, "fpr": 0.336538, "n": 539, "positive_rate": 0.640074, "tpr": 0.830816}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "JungleLee/bert-toxic-comment-classification", "num_labels": 2, "revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "task_type": "binary_classification"}
model_ref	JungleLee/bert-toxic-comment-classification
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	9c59104291a49a877078424cbfd9f3af23c26ef9
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.273143
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.216557
demographic_parity_difference	0.453158

Evidence:

Evidence ID	Hash	URI
06a63a53-503e-407e-9fd5-bb7541a78f95	sha256:ab70f65c4a65f22c3135bb29025cb250c6f39a14a3027ea26a913085fd3cb051	s3://trustlens/evidence/5f5c23d1-496e-4ccb-ae94-67ae9624ccd6/06a63a53-503e-407e-9fd5-bb7541a78f95.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign 0/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	JungleLee/bert-toxic-comment-classification
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.025333, "ci_upper": 0.046, "method": "bootstrap_percentile", "point": 0.036333}, "clean_accuracy": {"ci_lower": 0.832986, "ci_upper": 0.858795, "method": "wilson_score", "point": 0.846333}, "robust_accuracy": {"ci_lower": 0.795569, "ci_upper": 0.823638, "method": "wilson_score", "point": 0.81}}
probe_status	EVALUATED
accuracy_drop	0.036333
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.8463
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	9c59104291a49a877078424cbfd9f3af23c26ef9
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.81
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9571
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.058

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.036333, ci_lower=0.025333, ϵ =0.05)
relative_degradation	0.9571
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
c4ccb8a9-b4fd-47b0-81c1-a5eb15b3a3a2	sha256:a9787971d003cad1e7b91d69f0cda32dbf627f2d9f05cb8a8ac5b600049883eb	s3://trustlens/evidence/5f5c23d1-496e-4ccb-ae94-67ae9624ccd6/c4ccb8a9-b4fd-47b0-81c1-a5eb15b3a3a2.json

INTEGRITY (confidence 1.0)

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "1076 chars", "pass": true}, "files_listed": {"detail": "7 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:eae72ff94410074a31e1732a4fcc628b40b5ab875e1f639edc2fe2add8b42942", "model_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "pass": true}, "license_declared": {"detail": "afl-3.0", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, evaluation", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:eae72ff94410074a31e1732a4fcc628b40b5ab875e1f639edc2fe2add8b42942", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", "README.md", "config.json", "pytorch_model.bin", "tokenizer.json", "tokenizer_config.json", "vocab.txt"], "live_files_fingerprint": "sha256:eae72ff94410074a31e1732a4fcc628b40b5ab875e1f639edc2fe2add8b42942", "revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf hub lfs", "value": "924aff2eb02a6562097e6b09c543b010175abef67534e342baf501428977cc73"}, "weight_hash": {"algo": "sha256", "value": "924aff2eb02a6562097e6b09c543b010175abef67534e342baf501428977cc73"}}</pre>
disclosure	<pre>{"card_chars": 1076, "card_present": true, "license_structured": "afl-3.0", "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	no_material_risk
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>[]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "pytorch_model.bin", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
c9abc78d-4eab-44b3-937f-e71ec1329e29	sha256:0405959aca4602935176ee9a1fbe10369455d60294d503560e2b9248d151be7a	s3://trustlens/evidence/5f5c23d1-496e-4ccb-ae94-67ae9624ccd6/c9abc78d-4eab-44b3-937f-e71ec1329e29.json

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_limitations

missing_ethical_considerations

needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	{"documentation_completeness": {"coverage_ratio": 0.4, "detail": "2/5 required sections present", "pass": false, "sections_present": 2, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "Evaluation results", "pass": true}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "missing or trivial body", "pass": false}, "training_data": {"detail": "Training data", "pass": true}}
claims	{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (2/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}
status	EVALUATED
sections	{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": "Evaluation results", "present": true}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": null, "present": false}, "training_data": {"matched_heading": "Training data", "present": true}}
card_chars	1076
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	{"failed_gates": [], "gates_passed": true}
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	{"architecture": {"matched_heading": "Model description", "present": true}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": "How to use", "present": true}}
claim_boundary	{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}
contradictions	[]
coverage_ratio	0.4
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	2
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
methodology_version	tl-explainability-v1.0
documentation_source	{"content_length": 1076, "documentation_content_hash": "sha256:e3457d4f3d696e61f1a8728f6c2215bc34367b990cddd96fc2219fab872964f0", "documentation_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/JungleLee/bert-toxic-comment-classification/blob/9c59104291a49a877078424cbfd9f3af23c26ef9/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "JungleLee/bert-toxic-comment-classification", "source_model_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9"}

Evidence:

Evidence ID	Hash	URI
	sha256:da81d5d1ca76a544d0fa396ac2cf144ec9eb3446d0f52f92361712acb144c538	

8b2535e6-
c871-42f7-
aa5c-
a7eeb44adefe

s3://trustlens/evidence/
5f5c23d1-496e-4ccb-
ae94-67ae9624ccd6/8b2535e6-
c871-42f7-aa5c-a7eeb44adefe.json

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<pre>{"data_disclosure": {"detail": "Training data", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</pre>
claims	<pre>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</pre>
status	EVALUATED
behavior	<pre>{"benign_fpr": 0.13875, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.5152925531914894, "harmful_recall": 0.7866666666666666, "min_severe_n": 30, "model_ref": "JungleLee/bert-toxic-comment-classification", "model_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "n_evaluated": 3000, "overall_fnr": 0.21333333333333335, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.1099290780141844, "severe_fnr_ci": [0.0785316489943205, 0.1518113101829985], "severe_n": 282, "status": "EVALUATED"}</pre>
fnr_ratio	0.5152925531914894
card_chars	1076
severe_fnr	0.1099290780141844
limitations	["Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned.", "aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
bonus_checks	<pre>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	risk_detected
checks_present	1
claim_boundary	<pre>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</pre>
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	<pre>{"data_disclosure": {"matched_heading": "Training data", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</pre>
risks_triggered	<pre>["S-GOV-DISCLOSURE-GAP"]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
high_impact_claims	<pre>[]</pre>
methodology_version	tl-safety-v1.0
documentation_source	<pre>{"content_length": 1076, "documentation_content hash": "sha256:e3457d4f3d696e61f1a8728f6c2215bc34367b990cdd96fc2219fab872964f0",</pre>

	<pre>"documentation_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9", "documentation_source_type": "model_card", "documentation_url": "https:// huggingface.co/JungleLee/bert-toxic-comment-classification/blob/ 9c59104291a49a877078424cbfd9f3af23c26ef9/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "JungleLee/bert-toxic-comment- classification", "source_model_revision": "9c59104291a49a877078424cbfd9f3af23c26ef9"}</pre>
--	--

Evidence:

Evidence ID	Hash	URI
ec664d2a-6884-42d2-978a-9e4c868fccf5	sha256:bccd4c8c5f912675150d1c2ce4286bf0fd60e38fa898b80a18435c790c15ab38	s3://trustlens/evidence/5f5c23d1-496e-4ccb-ae94ec664d2a-6884-42d2-978a

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.036333, ci_lower=0.025333, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559

INTEGRITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.40
Human O/S/D	O=4 · S=4 · D=4
FRIES dimension score	4.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	8	8	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.217. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.846, robust_accuracy=0.810, accuracy_drop=0.036. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=1.00. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	4	4	4	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.40. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.110, fnr_ratio=0.52. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	JungleLee/bert-toxic-comment-classification
Model revision	9c59104291a49a877078424cbfd9f3af23c26ef9
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/JungleLee/bert-toxic-comment-classification/blob/9c59104291a49a877078424cbfd9f3af23c26ef9/README.md	9c59104291	sha256:e3457d4f3d6	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:23.834742Z · canonical JSON is the source of truth; this PDF is a rendered projection.