

Model	gravitee-io/bert-tiny-toxicity
Evaluation ID	9172d105-fec0-46b9-a717-00e7cf19586f
Status	FINALIZED
Created	2026-10-02T16:19:55.754068Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:30.276354Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for gravitee-io/bert-tiny-toxicity: 6.5634/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.5634/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_ethical_considerations, missing_evaluation, missing_intended_use, missing_limitations, missing_misuse_risks, missing_privacy, missing_security_warnings, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

6.5634 / 10

Dimension	Score
SAFETY	7.5595
FAIRNESS	8.0
INTEGRITY	8.6535
ROBUSTNESS	6.6039
EXPLAINABILITY	2.0

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	gravitee-io/bert-tiny-toxicity
Frozen revision	8377eb9a344d74b9d12e656dc3f3632b74f33673

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	8	8	8
ROBUSTNESS	9	4	8
INTEGRITY	9	9	8
EXPLAINABILITY	2	2	2
SAFETY	9	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)**evaluation_contract_v2** **user_defined_local_dataset**

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.274791, "ci_upper": 0.359233, "method": "bootstrap_percentile", "point": 0.317469}
eo_ci	{"B": 1000, "ci_lower": 0.133171, "ci_upper": 0.27036, "method": "bootstrap_percentile", "point": 0.202976}
f1_ci	{"B": 1000, "ci_lower": 0.37556, "ci_upper": 0.474596, "method": "bootstrap_percentile", "point": 0.422069}
groups	{"0": {"f1": 0.330379, "fpr": 0.349909, "n": 2461, "positive_rate": 0.394961, "tpr": 0.762082}, "1": {"f1": 0.752448, "fpr": 0.552885, "n": 539, "positive_rate": 0.71243, "tpr": 0.812689}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf", "model_ref": "gravitee-io/bert-tiny-toxicity", "num_labels": 2, "revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673", "task_type": "binary_classification"}
model_ref	gravitee-io/bert-tiny-toxicity
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	8377eb9a344d74b9d12e656dc3f3632b74f33673
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.422069
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.202976
demographic_parity_difference	0.317469

Evidence:

Evidence ID	Hash	URI
e8d22a3f-5fa9-4c37-b276-a4ff813b6bcd	sha256:3bab14c0c921223a7f474f1d7d2a2e917a8fb04de6ca4c0f096bb1b1a41757ed	s3://trustlens/evidence/9172d105-fec0-46b9-a717-00e7cf19586f/e8d22a3f-5fa9-4c37-b276-a4ff813b6bcd.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	gravitee-io/bert-tiny-toxicity
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.041333, "ci_upper": 0.068667, "method": "bootstrap_percentile", "point": 0.056333}, "clean_accuracy": {"ci_lower": 0.646897, "ci_upper": 0.680683, "method": "wilson_score", "point": 0.664}, "robust_accuracy": {"ci_lower": 0.590067, "ci_upper": 0.624991, "method": "wilson_score", "point": 0.607667}}
probe_status	EVALUATED
accuracy_drop	0.056333
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.664
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	8377eb9a344d74b9d12e656dc3f3632b74f33673
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.6077
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9152
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.095

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.056333, ci_lower=0.041333, ε=0.05)
relative_degradation	0.9152
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
0fd029b6-60ea-4f4b-929d-72b33b69b654	sha256:3b88fb2f115df36f578082814991fcc99f7b4090170ca9729796b5f7ef5ca4f0	s3://trustlens/evidence/fec0-46b9-a717-00e7cf190fd029b6-60ea-4f4b-929d

INTEGRITY (confidence 1.0)

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "6500 chars", "pass": true}, "files_listed": {"detail": "13 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:667b9f317afe08617f21eabf64c6459cd84f354bf686441a74f2f291865ddf4b", "model_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673", "pass": true}, "license_declared": {"detail": "openrail++", "pass": true}, "reproducibility_claims": {"detail": "signal groups: training data, evaluation", "pass": true}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:667b9f317afe08617f21eabf64c6459cd84f354bf686441a74f2f291865ddf4b", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", ".gitignore", "LICENSE", "NOTICE", "README.md", "config.json", "model.onnx", "model.quant.onnx", "model.safetensors", "special_tokens_map.json", "tokenizer.json", "tokenizer_config.json", "vocab.txt"], "live_files_fingerprint": "sha256:667b9f317afe08617f21eabf64c6459cd84f354bf686441a74f2f291865ddf4b", "revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf_hub_lfs", "value": "b9adbb07ffc75147b128e0d9177a4275c9dd864719c7a1335040021c7e333076"}, "weight_hash": {"algo": "sha256", "value": "b9adbb07ffc75147b128e0d9177a4275c9dd864719c7a1335040021c7e333076"}}</pre>
disclosure	<pre>{"card_chars": 6500, "card_present": true, "license_structured": "openrail++", "repro_groups_found": ["training_data", "evaluation"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	no_material_risk
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>[]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "model.safetensors", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
2373f7b3-3631-4aca-b417-b0745386a5e3	sha256:2ba1bd2b73de06aaa3c1b6d996f38656f54aa2aa66481541a594cf4f6a141fa2	s3://trustlens/evidence/9172d105-fec0-46b9-a717-00e7cf19586f/2373f7b3-3631-4aca-b417-b0745386a5e3.json

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_limitations

missing_evaluation

missing_ethical_considerations

needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	<code>{"documentation_completeness": {"coverage_ratio": 0.2, "detail": "1/5 required sections present", "pass": false, "sections_present": 1, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "missing or trivial body", "pass": false}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "missing or trivial body", "pass": false}, "training_data": {"detail": "card_data.datasets", "pass": true}}</code>
claims	<code>{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (1/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</code>
status	EVALUATED
sections	<code>{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": null, "present": false}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": null, "present": false}, "training_data": {"matched_heading": "card_data.datasets", "present": true}}</code>
card_chars	6500
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk_detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	<code>{"architecture": {"matched_heading": "card_data.base_model", "present": true}, "citation": {"matched_heading": "Citation", "present": true}, "examples": {"matched_heading": null, "present": false}}</code>
claim_boundary	<code>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</code>
contradictions	[]
coverage_ratio	0.2
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	1
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
methodology_version	tl-explainability-v1.0
documentation_source	<code>{"content_length": 6506, "documentation_content_hash": "sha256:87901062796d68a655122258d4a5b8ea15d0e591e5aebaa60aad476b550eb6b", "documentation_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/gravitee-io/bert-tiny-toxicity/blob/8377eb9a344d74b9d12e656dc3f3632b74f33673/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "gravitee-io/bert-tiny-toxicity", "source_model_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673"}</code>

Evidence:

Evidence ID	Hash	URI
d6caa6a8-3bc6-4bb3-bbf2-991c78ebb9ba	sha256:c03ac918083916155583e87b506299bb18b643239ac96896e7981511d317665a	s3://trustlens/evidence/9172d105-fec0-46b9-

	a717-00e7cf19586f/ d6caa6a8-3bc6-4bb3- bbf2-991c78ebb9ba.json
--	---

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "card_data.datasets", "pass": true}, "disclosure_completeness": {"checks_present": 1, "checks_required": 4, "coverage_ratio": 0.25, "detail": "1/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (1/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.3675, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.5403579871664979, "harmful_recall": 0.79, "min_severe_n": 30, "model_ref": "gravitee-io/bert-tiny-toxicity", "model_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673", "n_evaluated": 3000, "overall_fnr": 0.21, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.11347517730496454, "severe_fnr_ci": [0.08153515364767691, 0.15580468771901007], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.5403579871664979
card_chars	6500
severe_fnr	0.11347517730496454
limitations	[<code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	<code>[]</code>
aspect_scoring	risk_detected
checks_present	1
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.25
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": "card_data.datasets", "present": true}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	<code>["S-GOV-DISCLOSURE-GAP"]</code>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	<code>[]</code>
methodology_version	tl-safety-v1.0
documentation_source	<code>{"content_length": 6506, "documentation_content hash": "sha256:87901062796d68a655122258d4a5b8ea15d0e591e5aebaaf60aad476b550eb6b",</code>

```
"documentation_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673",
"documentation_source_type": "model_card", "documentation_url": "https://
huggingface.co/gravitee-io/bert-tiny-toxicity/blob/
8377eb9a344d74b9d12e656dc3f3632b74f33673/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "gravitee-io/bert-tiny-toxicity",
"source_model_revision": "8377eb9a344d74b9d12e656dc3f3632b74f33673"}
```

Evidence:

Evidence ID	Hash	URI
fbdd049f-3f91-4771-8e6a-3d02fa6ab708	sha256:49ffef59ffee291a3326fe10038a1f6080c49630836dfd1b05f0c267238278b7	s3://trustlens/evidence/fec0-46b9-a717-00e7cf19fbdd049f-3f91-4771-8e6a

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.056333, ci_lower=0.041333, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=4 · D=8
FRIES dimension score	6.604

INTEGRITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.20
Human O/S/D	O=2 · S=2 · D=2
FRIES dimension score	2.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.25
Human O/S/D	O=9 · S=6 · D=8
FRIES dimension score	7.559
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	8	8	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.203. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	4	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.664, robust_accuracy=0.608, accuracy_drop=0.056. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=1.00. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	2	2	2	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.20. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	9	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.113, fnr_ratio=0.54. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	gravitee-io/bert-tiny-toxicity
Model revision	8377eb9a344d74b9d12e656dc3f3632b74f33673
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/gravitee-io/bert-tiny-toxicity/blob/8377eb9a344d74b9d12e656dc3f3632b74f33673/README.md	8377eb9a34	sha256:87901062796	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:30.276354Z · canonical JSON is the source of truth; this PDF is a rendered projection.