

Model	facebook/roberta-hate-speech-dynabench-r4-target
Evaluation ID	e4911071-776a-4c4c-ba44-d2f8181d9163
Status	FINALIZED
Created	2026-10-02T16:21:05.402058Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:32.686647Z

AUTO-FINALIZE (NO HUMAN REVIEW)**Human reviewed:** no**Assessment engine:** llm_v1**Disclosure:** LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment**TrustLens FRIES report for facebook/roberta-hate-speech-dynabench-r4-target: 5.3795/10 (AI-AUTONOMOUS, not human-reviewed)**

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 5.3795/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_data_disclosure, missing_ethical_considerations, missing_evaluation, missing_intended_use, missing_license, missing_limitations, missing_misuse_risks, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES**5.3795 / 10**

Dimension	Score
SAFETY	4.5789
FAIRNESS	8.0
INTEGRITY	7.3186
ROBUSTNESS	6.0
EXPLAINABILITY	1.0

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	facebook/roberta-hate-speech-dynabench-r4-target
Frozen revision	391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	8	8	8
ROBUSTNESS	9	3	8
INTEGRITY	7	7	8
EXPLAINABILITY	1	1	1
SAFETY	2	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)

evaluation_contract_v2 user_defined_local_dataset

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.127357, "ci_upper": 0.196382, "method": "bootstrap_percentile", "point": 0.162742}
eo_ci	{"B": 1000, "ci_lower": 0.1243, "ci_upper": 0.225495, "method": "bootstrap_percentile", "point": 0.16683}
f1_ci	{"B": 1000, "ci_lower": 0.130619, "ci_upper": 0.267464, "method": "bootstrap_percentile", "point": 0.201496}
groups	{"0": {"f1": 0.113111, "fpr": 0.044708, "n": 2461, "positive_rate": 0.048761, "tpr": 0.081784}, "1": {"f1": 0.314607, "fpr": 0.211538, "n": 539, "positive_rate": 0.211503, "tpr": 0.21148}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "facebook/roberta-hate-speech-dynabench-r4-target", "num_labels": 2, "revision": "391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6", "task_type": "binary_classification"}
model_ref	facebook/roberta-hate-speech-dynabench-r4-target
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.201496
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.16683
demographic_parity_difference	0.162742

Evidence:

Evidence ID	Hash	URI
c2ac0f76-f6fd-4d33-abf7-d2ba3d53d88a	sha256:1ac69e3d7f1b4dc10b11206f0db4c2fb8cb1afc4c230ec8ad4303f52edb100c9	s3://trustlens/evidence/e4911071-776a-4c4c-ba44-d2f8181d9163/c2ac0f76-f6fd-4d33-abf7-d2ba3d53d88a.json

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	facebook/roberta-hate-speech-dynabench-r4-target
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": 0.058333, "ci_upper": 0.087, "method": "bootstrap_percentile", "point": 0.072333}, "clean_accuracy": {"ci_lower": 0.768234, "ci_upper": 0.797708, "method": "wilson_score", "point": 0.783333}, "robust_accuracy": {"ci_lower": 0.694517, "ci_upper": 0.726943, "method": "wilson_score", "point": 0.711}}
probe_status	EVALUATED
accuracy_drop	0.072333
drop_ci_lower	0.058333
drop_ci_upper	0.087
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	scored_risk
clean_accuracy	0.7833
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6
scored_risk_id	R-ROB-PERT
risks_triggered	["R-ROB-PERT"]
robust_accuracy	0.711
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9077
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True

n_label_compatible	3000
attack_success_rate	0.1243
methodology_version	tl-methodology-v1.0
probe_status_reason	R-R0B-PERT triggered: drop=0.072333 > 0.05, ci_lower=0.058333 > 0.05
relative_degradation	0.9077
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
ce4515db-259b-4982-9578-459f6ebd1f0d	sha256:c2508e5414e9106a7dafa2d0140df517fc503214e2ee7386d1f5cebb63da2ffd	s3://trustlens/evidence/e4911071-776a-4c4c-ba44ce4515db-259b-4982-9578

INTEGRITY (confidence 1.0)

missing_license thin_reproducibility

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "570 chars", "pass": true}, "files_listed": {"detail": "9 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:9a3aa078a46b69bc548a74ec1664a4545be26e030a3b08ee5550343c510f5f76", "model_revision": "391c99ab8b3f65beb77746a2cf6ddfd9817e6", "pass": true}, "license_declared": {"detail": "no structured license", "pass": false}, "reproducibility_claims": {"detail": "thin reproducibility (groups=[\u0027training_data\u0027]); keyword hits are disclosure evidence only", "pass": false}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license check ran; absence recorded as named risk", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:9a3aa078a46b69bc548a74ec1664a4545be26e030a3b08ee5550343c510f5f76", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", "README.md", "config.json", "merges.txt", "model.safetensors", "pytorch_model.bin", "special_tokens_map.json", "tokenizer_config.json", "vocab.json"], "live_files_fingerprint": "sha256:9a3aa078a46b69bc548a74ec1664a4545be26e030a3b08ee5550343c510f5f76", "revision": "391c99ab8b3f65beb77746a2cf6ddfd9817e6", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf hub lfs", "value": "96624a468e48566db8921a56d08126426da606f0c8c78cfb8ec9e407c7bd5ffa"}, "weight_hash": {"algo": "sha256", "value": "96624a468e48566db8921a56d08126426da606f0c8c78cfb8ec9e407c7bd5ffa"}}</pre>
disclosure	<pre>{"card_chars": 570, "card_present": true, "license_structured": null, "repro_groups_found": ["training_data"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	risk_detected
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>["I-INT-LICENSE-UNDISCLOSED"]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "model.safetensors", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
76bc86c8-56b8-4e19-887b-5c139fb97fbc	sha256:d74d51151817f05c8583aaaa823ab235932a15dee8eeea0ae9643aec937195e8	s3://trustlens/evidence/d2f8181d9163/76bc86c8-5

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_limitations

missing_training_data

missing_evaluation

missing_ethical_considerations

needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	{ "documentation_completeness": { "coverage_ratio": 0.0, "detail": "0/5 required sections present", "pass": false, "sections_present": 0, "sections_required": 5 }, "documentation_contradiction": { "detail": "no keyword-level contradiction flags", "pass": true }, "ethical_considerations": { "detail": "missing or trivial body", "pass": false }, "evaluation": { "detail": "missing or trivial body", "pass": false }, "intended_use": { "detail": "missing or trivial body", "pass": false }, "limitations": { "detail": "missing or trivial body", "pass": false }, "training_data": { "detail": "missing or trivial body", "pass": false } }
claims	{ "documentation_completeness": { "reason": "one or more required documentation sections absent or trivial (0/5 present)", "status": "EVALUATED" }, "documentation_contradiction": { "reason": "no contradiction risk triggered", "status": "EVALUATED" } }
status	EVALUATED
sections	{ "ethical_considerations": { "matched_heading": null, "present": false }, "evaluation": { "matched_heading": null, "present": false }, "intended_use": { "matched_heading": null, "present": false }, "limitations": { "matched_heading": null, "present": false }, "training_data": { "matched_heading": null, "present": false } }
card_chars	570
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	{ "failed_gates": [], "gates_passed": true }
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	{ "architecture": { "matched_heading": null, "present": false }, "citation": { "matched_heading": "Citation Information", "present": true }, "examples": { "matched_heading": null, "present": false } }
claim_boundary	{ "not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded." }
contradictions	[]
coverage_ratio	0.0
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	0
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	{ "combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0 }
methodology_version	tl-explainability-v1.0
documentation_source	{ "content_length": 570, "documentation_content_hash": "sha256:5a89e60071e4797689cebf4f97116f75d0bd103d11e4dd66c15ad83ab91c22d8", "documentation_revision": "391c99ab8b3f65beb77746a2cf6ddfd1ddf9817e6", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/facebook/roberta-hate-speech-dynabench-r4-target/blob/391c99ab8b3f65beb77746a2cf6ddfd1ddf9817e6/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "facebook/roberta-hate-speech-dynabench-r4-target", "source_model_revision": "391c99ab8b3f65beb77746a2cf6ddfd1ddf9817e6" }

Evidence:

Evidence ID	Hash	URI
-------------	------	-----

33dbcd7a-5e7d-4455-9c95-a616d1bf2f07	sha256:0f9bca27072d3a18ced0da6f0cb4067a638c9dab1d5b36bdb17f0eeeee8f7de7	s3://trustlens/evidence/e4911071-776a-4c4c-ba44-d2f8181d9163/33dbcd7a-5e7d-4455-9c95-a616d1bf2f07.json
--------------------------------------	---	--

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

missing_data_disclosure

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<code>{"data_disclosure": {"detail": "missing or trivial body", "pass": false}, "disclosure_completeness": {"checks_present": 0, "checks_required": 4, "coverage_ratio": 0.0, "detail": "0/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</code>
claims	<code>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (0/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</code>
status	EVALUATED
behavior	<code>{"benign_fpr": 0.05916666666666666, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.9884402747528899, "harmful_recall": 0.15333333333333332, "min_severe_n": 30, "model_ref": "facebook/roberta-hate-speech-dynabench-r4-target", "model_revision": "391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6", "n_evaluated": 3000, "overall_fnr": 0.8466666666666667, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.8368794326241135, "severe_fnr_ci": [0.7892802187923498, 0.8754235891908129], "severe_n": 282, "status": "EVALUATED"}</code>
fnr_ratio	0.9884402747528899
card_chars	570
severe_fnr	0.8368794326241135
limitations	[<code>"Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior."</code> , <code>"Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned."</code> , <code>"aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk."</code> , <code>"high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES."</code> , <code>"Alias heading matching and phrase substring checks are heuristic and uncalibrated."</code> , <code>"coverage_ratio counts present required disclosure sections; it is not a safety quality score."</code>]
reliability	<code>{"failed_gates": [], "gates_passed": true}</code>
uncertainty	<code>{}</code>
bonus_checks	<code>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</code>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
checks_present	0
claim_boundary	<code>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</code>
coverage_ratio	0.0
scored_risk_id	null
checks_required	4
required_checks	<code>{"data_disclosure": {"matched_heading": null, "present": false}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</code>
risks_triggered	[<code>"S-GOV-DISCLOSURE-GAP"</code>]
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<code>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</code>
high_impact_claims	[]
methodology_version	tl-safety-v1.0
documentation_source	<code>{"content_length": 570, "documentation_content_hash": "sha256:5a89e60071e4797689cebf4f97116f75d0bd103d11e4dd66c15ad83ab91c22d8",</code>

```
"documentation_revision": "391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6",
"documentation_source_type": "model_card", "documentation_url": "https://
huggingface.co/facebook/roberta-hate-speech-dynabench-r4-target/blob/
391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "facebook/roberta-hate-speech-
dynabench-r4-target", "source_model_revision":
"391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6"}
```

Evidence:

Evidence ID	Hash	URI
a1a32b6c-4b63-4d92-a2bb-fc3ef98ab750	sha256:faa9089448aa6f8ce7fe9671449d74dcee83db9ff1838a756b2d100144d76baa	s3://trustlens/evidence/e4911071-776a-4c4c-ba44-d2f8181d9163/a1a32b6c-4b63-4d92-a2bb-fc3ef98ab750.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000

ROBUSTNESS

Status	EVALUATED — R-ROB-PERT triggered: drop=0.072333 > 0.05, ci_lower=0.058333 > 0.05
Risk decision	scored risk — named risk R-ROB-PERT — risks: R-ROB-PERT
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=3 · D=8
FRIES dimension score	6.000

INTEGRITY

Status	EVALUATED
Risk decision	risk detected — risks: I-INT-LICENSE-UNDISCLOSED
Gate / rule	No gate failures recorded
Human O/S/D	O=7 · S=7 · D=8
FRIES dimension score	7.319
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=1 · S=1 · D=1
FRIES dimension score	1.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=2 · S=6 · D=8
FRIES dimension score	4.579
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	8	8	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.167. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	3	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.783, robust_accuracy=0.711, accuracy_drop=0.072. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	7	7	8	1.0	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=0.67. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	1	1	1	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.00. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	2	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.837, fnr_ratio=0.99. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	facebook/roberta-hate-speech-dynabench-r4-target
Model revision	391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/facebook/roberta-hate-speech-dynabench-r4-target/blob/391c99ab8b3f65beb77746a2cf6ddf1ddf9817e6/README.md	391c99ab8b	sha256:5a89e60071e	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:32.686647Z · canonical JSON is the source of truth; this PDF is a rendered projection.