

Model	Hate-speech-CNERG/dehatebert-mono-english
Evaluation ID	2c32080a-7586-4ebf-9c66-5801bf760ea6
Status	FINALIZED
Created	2026-10-02T16:23:24.829670Z
Report	report_v1 · version 1 · generated 2026-10-02T17:54:33.639324Z

AUTO-FINALIZE (NO HUMAN REVIEW)

Human reviewed: no

Assessment engine: llm_v1

Disclosure: LLM-assisted O/S/D: FAIRNESS/ROBUSTNESS come from the heuristic mapping; INTEGRITY/EXPLAINABILITY/SAFETY were proposed by an LLM (provider and model recorded per aspect) or by the heuristic when every LLM provider failed. Not human-reviewed, not ground truth.

Methodology status: LEGACY_HEURISTIC_OSD_V1. O = Occurrence, S = Severity, D = Detection (FRIES convention; higher is better/safer).

Executive Assessment

TrustLens FRIES report for Hate-speech-CNERG/dehatebert-mono-english: 6.1772/10 (AI-AUTONOMOUS, not human-reviewed)

- Evaluation mode: AI-AUTONOMOUS (workflow — not LLM interpretation)
- Original FRIES score: 6.1772/10 (not FRIES2)
- Human reviewed: no
- Probe flags: evaluation_contract_v2, missing_data_disclosure, missing_ethical_considerations, missing_evaluation, missing_intended_use, missing_limitations, missing_misuse_risks, missing_privacy, ...
- O/S/D are LLM-assisted proposals for INTEGRITY/EXPLAINABILITY/SAFETY (heuristic for FAIRNESS/ROBUSTNESS) — not ground truth.

FRIES

6.1772 / 10

Dimension	Score
SAFETY	4.5789
FAIRNESS	8.6535
INTEGRITY	8.0
ROBUSTNESS	8.6535
EXPLAINABILITY	1.0

Score type: original_FRIES · overall confidence 0.9966

Original FRIES computed from finalized O/S/D — not FRIES2. O/S/D on this path are LLM-assisted proposals (heuristic for FAIRNESS/ROBUSTNESS), not ground truth.

Evaluation Contract

Contract kind	documentation_only
Task	—
Dataset (comparability label)	—
Config	—
Pairing ID	—

Model Identity

Model reference	Hate-speech-CNERG/dehatebert-mono-english
Frozen revision	25d0e4d9122d2a5c283e07405a325e3dfd4a73b3

Dataset Identity

Dataset key	—
Dataset revision	—
User dataset ID	—
Dataset content hash	—

Local Execution Environment

Execution device	GPU (NVIDIA GeForce RTX 4060 Laptop GPU)
Inference backend	local_hf
CUDA available	True

Human O/S/D Assessment

O, S, and D are always human-entered — never inferred from probe evidence, never generated by a probe.

Aspect	O (Occurrence)	S (Severity)	D (Detection)
FAIRNESS	9	9	8
ROBUSTNESS	9	9	8
INTEGRITY	8	8	8
EXPLAINABILITY	1	1	1
SAFETY	2	6	8

Source: osd_agent_autonomous · human reviewed: no

Confidence summary

Method: geometric_mean_v1 · proposed calibration: True

Scope	Confidence
Overall	0.9966
FAIRNESS	0.983
ROBUSTNESS	1.0
INTEGRITY	1.0
EXPLAINABILITY	1.0
SAFETY	1.0

Evidence strength only — not correctness or O/S/D

FAIRNESS (confidence 0.983)**evaluation_contract_v2** **user_defined_local_dataset**

Metric	Value
note	Metrics are objective evidence only – not a normative fair/unfair judgment, not 0/5/D, and not product FRIES
seed	42
dp_ci	{"B": 1000, "ci_lower": 0.089156, "ci_upper": 0.151984, "method": "bootstrap_percentile", "point": 0.120086}
eo_ci	{"B": 1000, "ci_lower": 0.073056, "ci_upper": 0.181223, "method": "bootstrap_percentile", "point": 0.127371}
f1_ci	{"B": 1000, "ci_lower": 0.134879, "ci_upper": 0.283196, "method": "bootstrap_percentile", "point": 0.210064}
groups	{"0": {"f1": 0.117647, "fpr": 0.030566, "n": 2461, "positive_rate": 0.035758, "tpr": 0.078067}, "1": {"f1": 0.327711, "fpr": 0.076923, "n": 539, "positive_rate": 0.155844, "tpr": 0.205438}}
inference	{"backend": "local hf", "batch_size": 8, "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "dtype": "torch.float32", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local hf", "model_ref": "Hate-speech-CNERG/dehatebert-mono-english", "num_labels": 2, "revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3", "task_type": "binary_classification"}
model_ref	Hate-speech-CNERG/dehatebert-mono-english
predictor	inference_backend
excess_dpd	0.0
dataset_key	null
min_group_n	30
n_evaluated	3000
text_column	text
group_column	identity_ref
missing_rows	{"missing_count": 0, "missing_reasons": {}}
probe_status	EVALUATED
excluded_rows	{"missing_group": 0, "missing_target": 0, "missing_text": 0, "unmapped_label": 0}
fairness_mode	user_defined_local
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
label_rate_gap	0.504795
model_revision	25d0e4d9122d2a5c283e07405a325e3dfd4a73b3
excluded_groups	[]
observed_groups	[{"count": 2461, "value": "0"}, {"count": 539, "value": "1"}]
dataset_revision	null
evaluation_class	user_dataset_v2
proposed_mapping	False
dataset_row_count	3000
confidence_factors	{"combined": 0.983, "data_quality": 1.0, "evidence_completeness": 0.95, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
needs_human_review	True
subgroup_f1_spread	0.210064
probe_status_reason	null
sensitive_attribute	identity_ref
min_group_n_observed	539
positive_label_index	1

equalized_odds_difference	0.127371
demographic_parity_difference	0.120086

Evidence:

Evidence ID	Hash	URI
e26c8445-2ef4-4f1f-835c-5ca553ce033f	sha256:ab106176eb01b9ea48368c875b23b25c23d297f8476925cb01c07ebb66c2d788	s3://trustlens/evidence/2c32080a-7586-4ebf-9c66-e26c8445-2ef4-4f1f-835c

ROBUSTNESS (confidence 1.0)

evaluation_contract_v2

user_defined_local_dataset

Metric	Value
norm	discrete
note	Layer A evidence only – R-ROB-PERT detection does not assign O/S/D
seed	42
attack	char_swap
status	EVALUATED
dataset	{"evaluation_domain": null, "logical_key": null, "resolution_source": "user_dataset_v2"}
epsilon	0.03
inference	{"backend": "local_hf", "cuda_available": true, "device": "cuda", "device_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "device_reason": "CUDA GPU detected and selected", "execution_device": "cuda", "fallback_reason": null, "gpu_available": true, "gpu_name": "NVIDIA GeForce RTX 4060 Laptop GPU", "inference_backend": "local_hf"}
model_ref	Hate-speech-CNERG/dehatebert-mono-english
n_samples	3000
task_type	null
dataset_key	null
limitations	[]
max_changes	3
n_evaluated	3000
n_requested	3000
reliability	{"failed_gates": [], "gates_passed": true}
text_column	text
uncertainty	{"accuracy_drop": {"B": 1000, "ci_lower": -0.004333, "ci_upper": 0.006667, "method": "bootstrap_percentile", "point": 0.001333}, "clean_accuracy": {"ci_lower": 0.787358, "ci_upper": 0.81587, "method": "wilson_score", "point": 0.802}, "robust_accuracy": {"ci_lower": 0.78599, "ci_upper": 0.814574, "method": "wilson_score", "point": 0.800667}}
probe_status	EVALUATED
accuracy_drop	0.001333
excluded_rows	{"missing_target": 0, "missing_text": 0, "unmapped_label": 0}
label_mapping	[{"dataset_value": "0", "model_label_index": 0}, {"dataset_value": "1", "model_label_index": 1}]
target_column	label
aspect_scoring	no_material_risk
clean_accuracy	0.802
dataset_format	csv
label_encoding	{"0": 0, "1": 1}
model_revision	25d0e4d9122d2a5c283e07405a325e3dfd4a73b3
scored_risk_id	null
risks_triggered	[]
robust_accuracy	0.8007
dataset_revision	null
evaluation_class	user_dataset_v2
n_perturb_failed	0
proposed_mapping	False
degradation_ratio	0.9983
confidence_factors	{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}
dataset_content_id	c5d8778e-1eba-4a8c-a55c-82487c90a1b2
inference_executed	True
n_label_compatible	3000
attack_success_rate	0.0117

methodology_version	tl-methodology-v1.0
probe_status_reason	no material perturbation risk at 95% CI (drop=0.001333, ci_lower=-0.004333, ϵ =0.05)
relative_degradation	0.9983
label_compat_fraction	1.0
perturbation_coverage	1.0
n_successfully_perturbed	3000

Evidence:

Evidence ID	Hash	URI
a180ea33-c8ae-4002-9861-758fd109530b	sha256:8c1841603b11d581754542ee0a25ada2d43be88ac2cd57879a59b67bd1b936d9	s3://trustlens/evidence/2c32080a-7586-4ebf-9c66-5801bf76a180ea33-c8ae-4002-9861-758fd109530b.json

INTEGRITY (confidence 1.0)

thin_reproducibility

Metric	Value
note	Layer A evidence only – Integrity risks do not assign O/S/D
checks	<pre>{"card_present": {"detail": "1046 chars", "pass": true}, "files_listed": {"detail": "8 files; config.json present; weight-like filename present", "pass": true}, "files_listing_recorded": {"detail": "recorded Hub file listing fingerprint (not weight bytes)", "files_listing_fingerprint": "sha256:d86be1852cdff748819bf551160d2632d6f144dc19be32d7c8c70e21fe69aeb8", "model_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3", "pass": true}, "license_declared": {"detail": "apache-2.0", "pass": true}, "reproducibility_claims": {"detail": "thin reproducibility (groups=[\u0027hyperparameters\u0027]); keyword hits are disclosure evidence only", "pass": false}, "revision_pinned": {"detail": "sha-like revision", "pass": true}}</pre>
claims	<pre>{"cryptographic_identity": {"reason": "local artifact hash matches trusted reference", "status": "EVALUATED"}, "files_listing_currency": {"reason": "Current Hub file listing matches the import-time recorded fingerprint.", "status": "EVALUATED"}, "identity_recording": {"reason": "identity metadata recorded from import snapshot", "status": "EVALUATED"}, "license_disclosure": {"reason": "structured license present", "status": "EVALUATED"}}</pre>
status	EVALUATED
identity	<pre>{"files_listing_fingerprint": "sha256:d86be1852cdff748819bf551160d2632d6f144dc19be32d7c8c70e21fe69aeb8", "files_listing_reverification": "match", "hash_comparison": "match", "hub_files": [".gitattributes", "README.md", "config.json", "flax_model.msgpack", "pytorch_model.bin", "special_tokens_map.json", "tokenizer_config.json", "vocab.txt"], "live_files_fingerprint": "sha256:d86be1852cdff748819bf551160d2632d6f144dc19be32d7c8c70e21fe69aeb8", "revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3", "sha_like": true, "trusted_reference": {"algo": "sha256", "provenance": null, "retrieved_at": null, "source": "hf_hub_lfs", "value": "ce02a5f74e55cff488c487f52773af12f908bbefd0a9267fc19c16906bdc0dd3"}, "weight_hash": {"algo": "sha256", "value": "ce02a5f74e55cff488c487f52773af12f908bbefd0a9267fc19c16906bdc0dd3"}}</pre>
disclosure	<pre>{"card_chars": 1046, "card_present": true, "license_structured": "apache-2.0", "repro_groups_found": ["hyperparameters"]}</pre>
limitations	["Integrity risks measure disclosure and identity recording, not tampering or legal compliance.", "Named Integrity risks are not additive FRIES occurrences and must not be summed into 0.", "Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.", "Hub file listing fingerprint identifies sorted filenames, not model weight bytes.", "SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.", "Model card text may reflect the Hub default branch, not the pinned revision.", "Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering \u2014 it never downloads or hashes weight files."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	no_material_risk
claim_boundary	<pre>{"cryptographic_identity": "local artifact hash matches trusted reference", "files_listing_currency": "Current Hub file listing matches the import-time recorded fingerprint."}</pre>
scored_risk_id	null
risks_triggered	<pre>[]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-integrity-v1.0
artifact_verification	<pre>{"file": "pytorch_model.bin", "performed": true}</pre>

Evidence:

Evidence ID	Hash	URI
d3cfa5d3-19ba-49b2-8c48-523b32b4c58f	sha256:81a2b4741fa7e1c3e2eef4e78127ef7a1380a47dace41dd95111e83961c56f6f	s3://trustlens/evidence/2c32080a-7586-4ebf-9c66-d3cfa5d3-19ba-49b2-8c48

EXPLAINABILITY (confidence 1.0)

missing_intended_use

missing_limitations

missing_training_data

missing_evaluation

missing_ethical_considerations

needs_human_review

Metric	Value
note	Layer A documentation/transparency evidence only – not O/S/D or interpretability quality
checks	<pre>{"documentation_completeness": {"coverage_ratio": 0.0, "detail": "0/5 required sections present", "pass": false, "sections_present": 0, "sections_required": 5}, "documentation_contradiction": {"detail": "no keyword-level contradiction flags", "pass": true}, "ethical_considerations": {"detail": "missing or trivial body", "pass": false}, "evaluation": {"detail": "missing or trivial body", "pass": false}, "intended_use": {"detail": "missing or trivial body", "pass": false}, "limitations": {"detail": "missing or trivial body", "pass": false}, "training_data": {"detail": "missing or trivial body", "pass": false}}</pre>
claims	<pre>{"documentation_completeness": {"reason": "one or more required documentation sections absent or trivial (0/5 present)", "status": "EVALUATED"}, "documentation_contradiction": {"reason": "no contradiction risk triggered", "status": "EVALUATED"}}</pre>
status	EVALUATED
sections	<pre>{"ethical_considerations": {"matched_heading": null, "present": false}, "evaluation": {"matched_heading": null, "present": false}, "intended_use": {"matched_heading": null, "present": false}, "limitations": {"matched_heading": null, "present": false}, "training_data": {"matched_heading": null, "present": false}}</pre>
card_chars	1046
limitations	["Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.", "Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.", "aspect_scoring=risk detected means a documentation risk was detected \u2014 not a FRIES-scored risk.", "Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.", "coverage_ratio counts present required sections; it is not an explainability quality score."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
probe_status	EVALUATED
osd_proposals	[]
aspect_scoring	risk_detected
bonus_sections	<pre>{"architecture": {"matched_heading": null, "present": false}, "citation": {"matched_heading": null, "present": false}, "examples": {"matched_heading": null, "present": false}}</pre>
claim_boundary	<pre>{"not_established": "This does not establish explanation faithfulness, causal correctness, human usefulness, SHAP/attention availability, or general explainability quality. Documentation coverage is not proof that the model is explainable.", "supported": "Required model-card documentation sections were checked for presence with non-trivial body text; coverage counts and keyword-level contradiction flags were recorded."}</pre>
contradictions	[]
coverage_ratio	0.0
scored_risk_id	null
risks_triggered	["E-DOC-INCOMPLETE"]
proposed_mapping	False
sections_present	0
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
sections_required	5
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
methodology_version	tl-explainability-v1.0
documentation_source	<pre>{"content_length": 1047, "documentation_content_hash": "sha256:4b08be134f8ca23e0c10b35c975a3923258d5aed7e25f4b6090c8009e7625a4a", "documentation_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3", "documentation_source_type": "model_card", "documentation_url": "https://huggingface.co/Hate-speech-CNERG/dehatebert-mono-english/blob/25d0e4d9122d2a5c283e07405a325e3dfd4a73b3/README.md", "retrieval_error": null, "retrieval_status": "ok", "source_model_ref": "Hate-speech-CNERG/dehatebert-mono-english", "source_model_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3"}</pre>

Evidence:

Evidence ID	Hash	URI
-------------	------	-----

14b6ca17-5980-4612-a849-8f2f55ef6b06	sha256:662ed2b6c95a3b48579eb4bd30fbf326a31be0415cb9d55fb22d8a56051ca667	s3://trustlens/evidence/2c32080a-7586-4ebf-9c66-5801bf760ea6/14b6ca17-5980-4612-a849-8f2f55ef6b06.json
--------------------------------------	---	--

SAFETY (confidence 1.0)

missing_misuse_risks

missing_privacy

missing_security_warnings

missing_data_disclosure

needs_human_review

Metric	Value
note	Layer A safety-governance disclosure evidence only – not O/S/D, not behavioral safety, not FRIES2 caps
checks	<pre>{"data_disclosure": {"detail": "missing or trivial body", "pass": false}, "disclosure_completeness": {"checks_present": 0, "checks_required": 4, "coverage_ratio": 0.0, "detail": "0/4 required sections present", "pass": false}, "misuse_risks": {"detail": "missing or trivial body", "pass": false}, "privacy": {"detail": "missing or trivial body", "pass": false}, "security_warnings": {"detail": "missing or trivial body", "pass": false}}</pre>
claims	<pre>{"disclosure_completeness": {"reason": "one or more required safety-disclosure sections absent or trivial (0/4 present)", "status": "EVALUATED"}, "phrase_matches": {"reason": "no listed high-impact phrases matched", "status": "EVALUATED"}}</pre>
status	EVALUATED
behavior	<pre>{"benign_fpr": 0.034583333333333334, "dataset_content_hash": "ef8ff36af2b476ffbad769eef1684b1850cf58e362da5a844fdc58c14078057d", "dataset_content_id": "c5d8778e-leba-4a8c-a55c-82487c90a1b2", "flags": [], "fnr_ratio": 0.9534912770121164, "harmful_recall": 0.14833333333333332, "min_severe_n": 30, "model_ref": "Hate-speech-CNERG/dehatebert-mono-english", "model_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3", "n_evaluated": 3000, "overall_fnr": 0.8516666666666667, "positive_label_index": 1, "severe_column": "severe", "severe_fnr": 0.8120567375886525, "severe_fnr_ci": [0.7623792982499114, 0.8533463344081762], "severe_n": 282, "status": "EVALUATED"}</pre>
fnr_ratio	0.9534912770121164
card_chars	1046
severe_fnr	0.8120567375886525
limitations	["Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.", "Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored risk_id remains null and no FRIES occurrence is assigned.", "aspect scoring=risk detected means a governance disclosure gap was detected \u2014 not a FRIES-scored risk.", "high_impact_claims are documentation metadata flags; they are never promoted to risks triggered and do not map to O/S/D or FRIES.", "Alias heading matching and phrase substring checks are heuristic and uncalibrated.", "coverage_ratio counts present required disclosure sections; it is not a safety quality score."]
reliability	<pre>{"failed_gates": [], "gates_passed": true}</pre>
uncertainty	<pre>{}</pre>
bonus_checks	<pre>{"bias_and_fairness_risks": {"matched_heading": null, "present": false}, "human_oversight": {"matched_heading": null, "present": false}}</pre>
probe_status	EVALUATED
osd_proposals	<pre>[]</pre>
aspect_scoring	risk_detected
checks_present	0
claim_boundary	<pre>{"not_established": "This does not establish model safety or unsafety, actual deployment risk, mitigation effectiveness, harm severity, policy compliance, medical/financial fitness, or refusal capability. Phrase matches (high impact claims) are lexical documentation evidence only \u2014 not evidence that the model is high-risk, unsafe, harmful, or safety-critical.", "supported": "Required safety-disclosure model-card sections were checked for presence with non-trivial body text; documentation completeness counts and lexical presence of listed high-impact phrases were recorded as documentation metadata flags."}</pre>
coverage_ratio	0.0
scored_risk_id	null
checks_required	4
required_checks	<pre>{"data_disclosure": {"matched_heading": null, "present": false}, "misuse_risks": {"matched_heading": null, "present": false}, "privacy": {"matched_heading": null, "present": false}, "security_warnings": {"matched_heading": null, "present": false}}</pre>
risks_triggered	<pre>["S-GOV-DISCLOSURE-GAP"]</pre>
proposed_mapping	False
methodology_basis	TRUSTLENS_FIVE_PROBE_METHODODOLOGY_AUDIT.md
confidence_factors	<pre>{"combined": 1.0, "data_quality": 1.0, "evidence_completeness": 1.0, "probe_reliability": 1.0}</pre>
high_impact_claims	<pre>[]</pre>
methodology_version	tl-safety-v1.0
documentation_source	<pre>{"content_length": 1047, "documentation_content_hash": "sha256:4b08be134f8ca23e0c10b35c975a3923258d5aed7e25f4b6090c8009e7625a4a",</pre>

```
"documentation_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3",
"documentation_source_type": "model card", "documentation_url": "https://
huggingface.co/Hate-speech-CNERG/dehatebert-mono-english/blob/
25d0e4d9122d2a5c283e07405a325e3dfd4a73b3/README.md", "retrieval_error": null,
"retrieval_status": "ok", "source_model_ref": "Hate-speech-CNERG/dehatebert-mono-
english", "source_model_revision": "25d0e4d9122d2a5c283e07405a325e3dfd4a73b3"}
```

Evidence:

Evidence ID	Hash	URI
5ca7a883-2bee-4f1a-b1de-13e2d21341c4	sha256:0dee1f34ca766cb0ff40a83e70d85997be628a0aa282a5e45449b69117819a64	s3://trustlens/evidence/2c32080a-7586-4ebf-9c66-5801bf760ea6/5ca7a883-2bee-4f1a-b1de-13e2d21341c4.json

Evidence Traceability

Per dimension: conclusion/risk-status → gate/rule → metric (see Dimension Evidence above, not repeated here) → evidence reference → human O/S/D → FRIES dimension score.

FAIRNESS

Status	EVALUATED
Risk decision	Not scored
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

ROBUSTNESS

Status	EVALUATED — no material perturbation risk at 95% CI (drop=0.001333, ci_lower=-0.004333, ε=0.05)
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=9 · S=9 · D=8
FRIES dimension score	8.653

INTEGRITY

Status	EVALUATED
Risk decision	no material risk
Gate / rule	No gate failures recorded
Human O/S/D	O=8 · S=8 · D=8
FRIES dimension score	8.000
Limitations	Integrity risks measure disclosure and identity recording, not tampering or legal compliance.; Named Integrity risks are not additive FRIES occurrences and must not be summed into O.; Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.; Hub file listing fingerprint identifies sorted filenames, not model weight bytes.; SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.; Model card text may reflect the Hub default branch, not the pinned revision.; Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.

EXPLAINABILITY

Status	EVALUATED
Risk decision	risk detected — risks: E-DOC-INCOMPLETE
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=1 · S=1 · D=1
FRIES dimension score	1.000
Limitations	Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.; Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.; Alias heading_matching and keyword contradiction checks are heuristic and uncalibrated.; coverage_ratio counts present required sections; it is not an explainability quality score.

SAFETY

Status	EVALUATED
Risk decision	risk detected — risks: S-GOV-DISCLOSURE-GAP
Gate / rule	No gate failures recorded
Documentation coverage	0.00
Human O/S/D	O=2 · S=6 · D=8
FRIES dimension score	4.579
Limitations	Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.; Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.; aspect_scoring=risk_detected means a governance disclosure gap was detected — not a FRIES-scored risk.; high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.; Alias heading matching and phrase substring checks are heuristic and uncalibrated.; coverage_ratio counts present required disclosure sections; it is not a safety quality score.

O/S/D representation

Methodology status: **LEGACY_HEURISTIC_OSD_V1** — LLM-assisted O/S/D suggestion — INTEGRITY/EXPLAINABILITY/SAFETY proposed by an LLM (heuristic fallback when all providers fail), FAIRNESS/ROBUSTNESS by the heuristic; proposed, not ground truth; requires human validation.

Aspect	O (Occurrence)	S (Severity)	D (Detection)	Evidence strength	Rationale
FAIRNESS	9	9	8	0.983	[PROPOSED / REQUIRES VALIDATION] FAIRNESS: observed disparity gap=0.127. Heuristic metric-to-O/S/D mapping — not validated science.
ROBUSTNESS	9	9	8	1.0	[PROPOSED / REQUIRES VALIDATION] ROBUSTNESS: clean_accuracy=0.802, robust_accuracy=0.801, accuracy_drop=0.001. Heuristic metric-to-O/S/D mapping — not validated science.
INTEGRITY	8	8	8	1.0	[PROPOSED / REQUIRES VALIDATION] INTEGRITY: metadata check pass rate=0.83. Heuristic metric-to-O/S/D mapping — not validated science.
EXPLAINABILITY	1	1	1	1.0	[PROPOSED / REQUIRES VALIDATION] EXPLAINABILITY: card coverage_ratio=0.00. Heuristic metric-to-O/S/D mapping — not validated science.
SAFETY	2	6	8	1.0	[PROPOSED / REQUIRES VALIDATION] SAFETY: severe_fnr=0.812, fnr_ratio=0.95. Heuristic metric-to-O/S/D mapping — not validated science.

OSD overall evidence-strength mean: 0.9966 (not FRIES certainty)

Reproducibility Information

Model reference	Hate-speech-CNERG/dehatebert-mono-english
Model revision	25d0e4d9122d2a5c283e07405a325e3dfd4a73b3
TrustLens version	—
Contract kind	documentation_only

Re-running this exact contract (same model + dataset revision) uses the same local inference path. This is reproducible configuration, not a claim of bitwise-identical output.

Documentation Sources

Source	Type	URL	Revision	Hash	Retrieval
HF card (pinned)	model_card	https://huggingface.co/Hate-speech-CNERG/dehatebert-mono-english/blob/25d0e4d9122d2a5c283e07405a325e3dfd4a73b3/README.md	25d0e4d912	sha256:4b08be134f8	ok

Limitations

- Alias heading matching and keyword contradiction checks are heuristic and uncalibrated.
- Alias heading matching and phrase substring checks are heuristic and uncalibrated.
- Explainability v1 measures model-card documentation completeness and simple consistency flags, not runtime explanations or attribution quality.
- Hub file listing fingerprint identifies sorted filenames, not model weight bytes.
- Integrity risks measure disclosure and identity recording, not tampering or legal compliance.
- Model card text may reflect the Hub default branch, not the pinned revision.
- Named E-DOC-* values are evidence-layer detections only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Named Integrity risks are not additive FRIES occurrences and must not be summed into O.
- Named S-GOV-DISCLOSURE-GAP is an evidence-layer detection only; scored_risk_id remains null and no FRIES occurrence is assigned.
- Operator-supplied trusted hashes are not a PKI; hash match does not prove publisher authenticity; mismatch does not prove tampering.
- Post-import file-listing re-verification detects a changed Hub listing, not weight-byte tampering — it never downloads or hashes weight files.
- SHA-like revision format is an uncalibrated heuristic, not cryptographic proof of immutability.
- Safety v1 measures model-card safety-disclosure completeness and lexical phrase matches, not runtime refusal or harm behavior.
- aspect_scoring=risk_detected means a documentation risk was detected — not a FRIES-scored risk.
- aspect_scoring=governance_detected means a governance disclosure gap was detected — not a FRIES-scored risk.
- coverage_ratio counts present required disclosure sections; it is not a safety quality score.
- coverage_ratio counts present required sections; it is not an explainability quality score.
- high_impact_claims are documentation metadata flags; they are never promoted to risks_triggered and do not map to O/S/D or FRIES.

report_v1 · report version 1 · generated 2026-10-02T17:54:33.639324Z · canonical JSON is the source of truth; this PDF is a rendered projection.